

22-10-22

CS60021: Scalable Data Minin

Streaming Algorithms

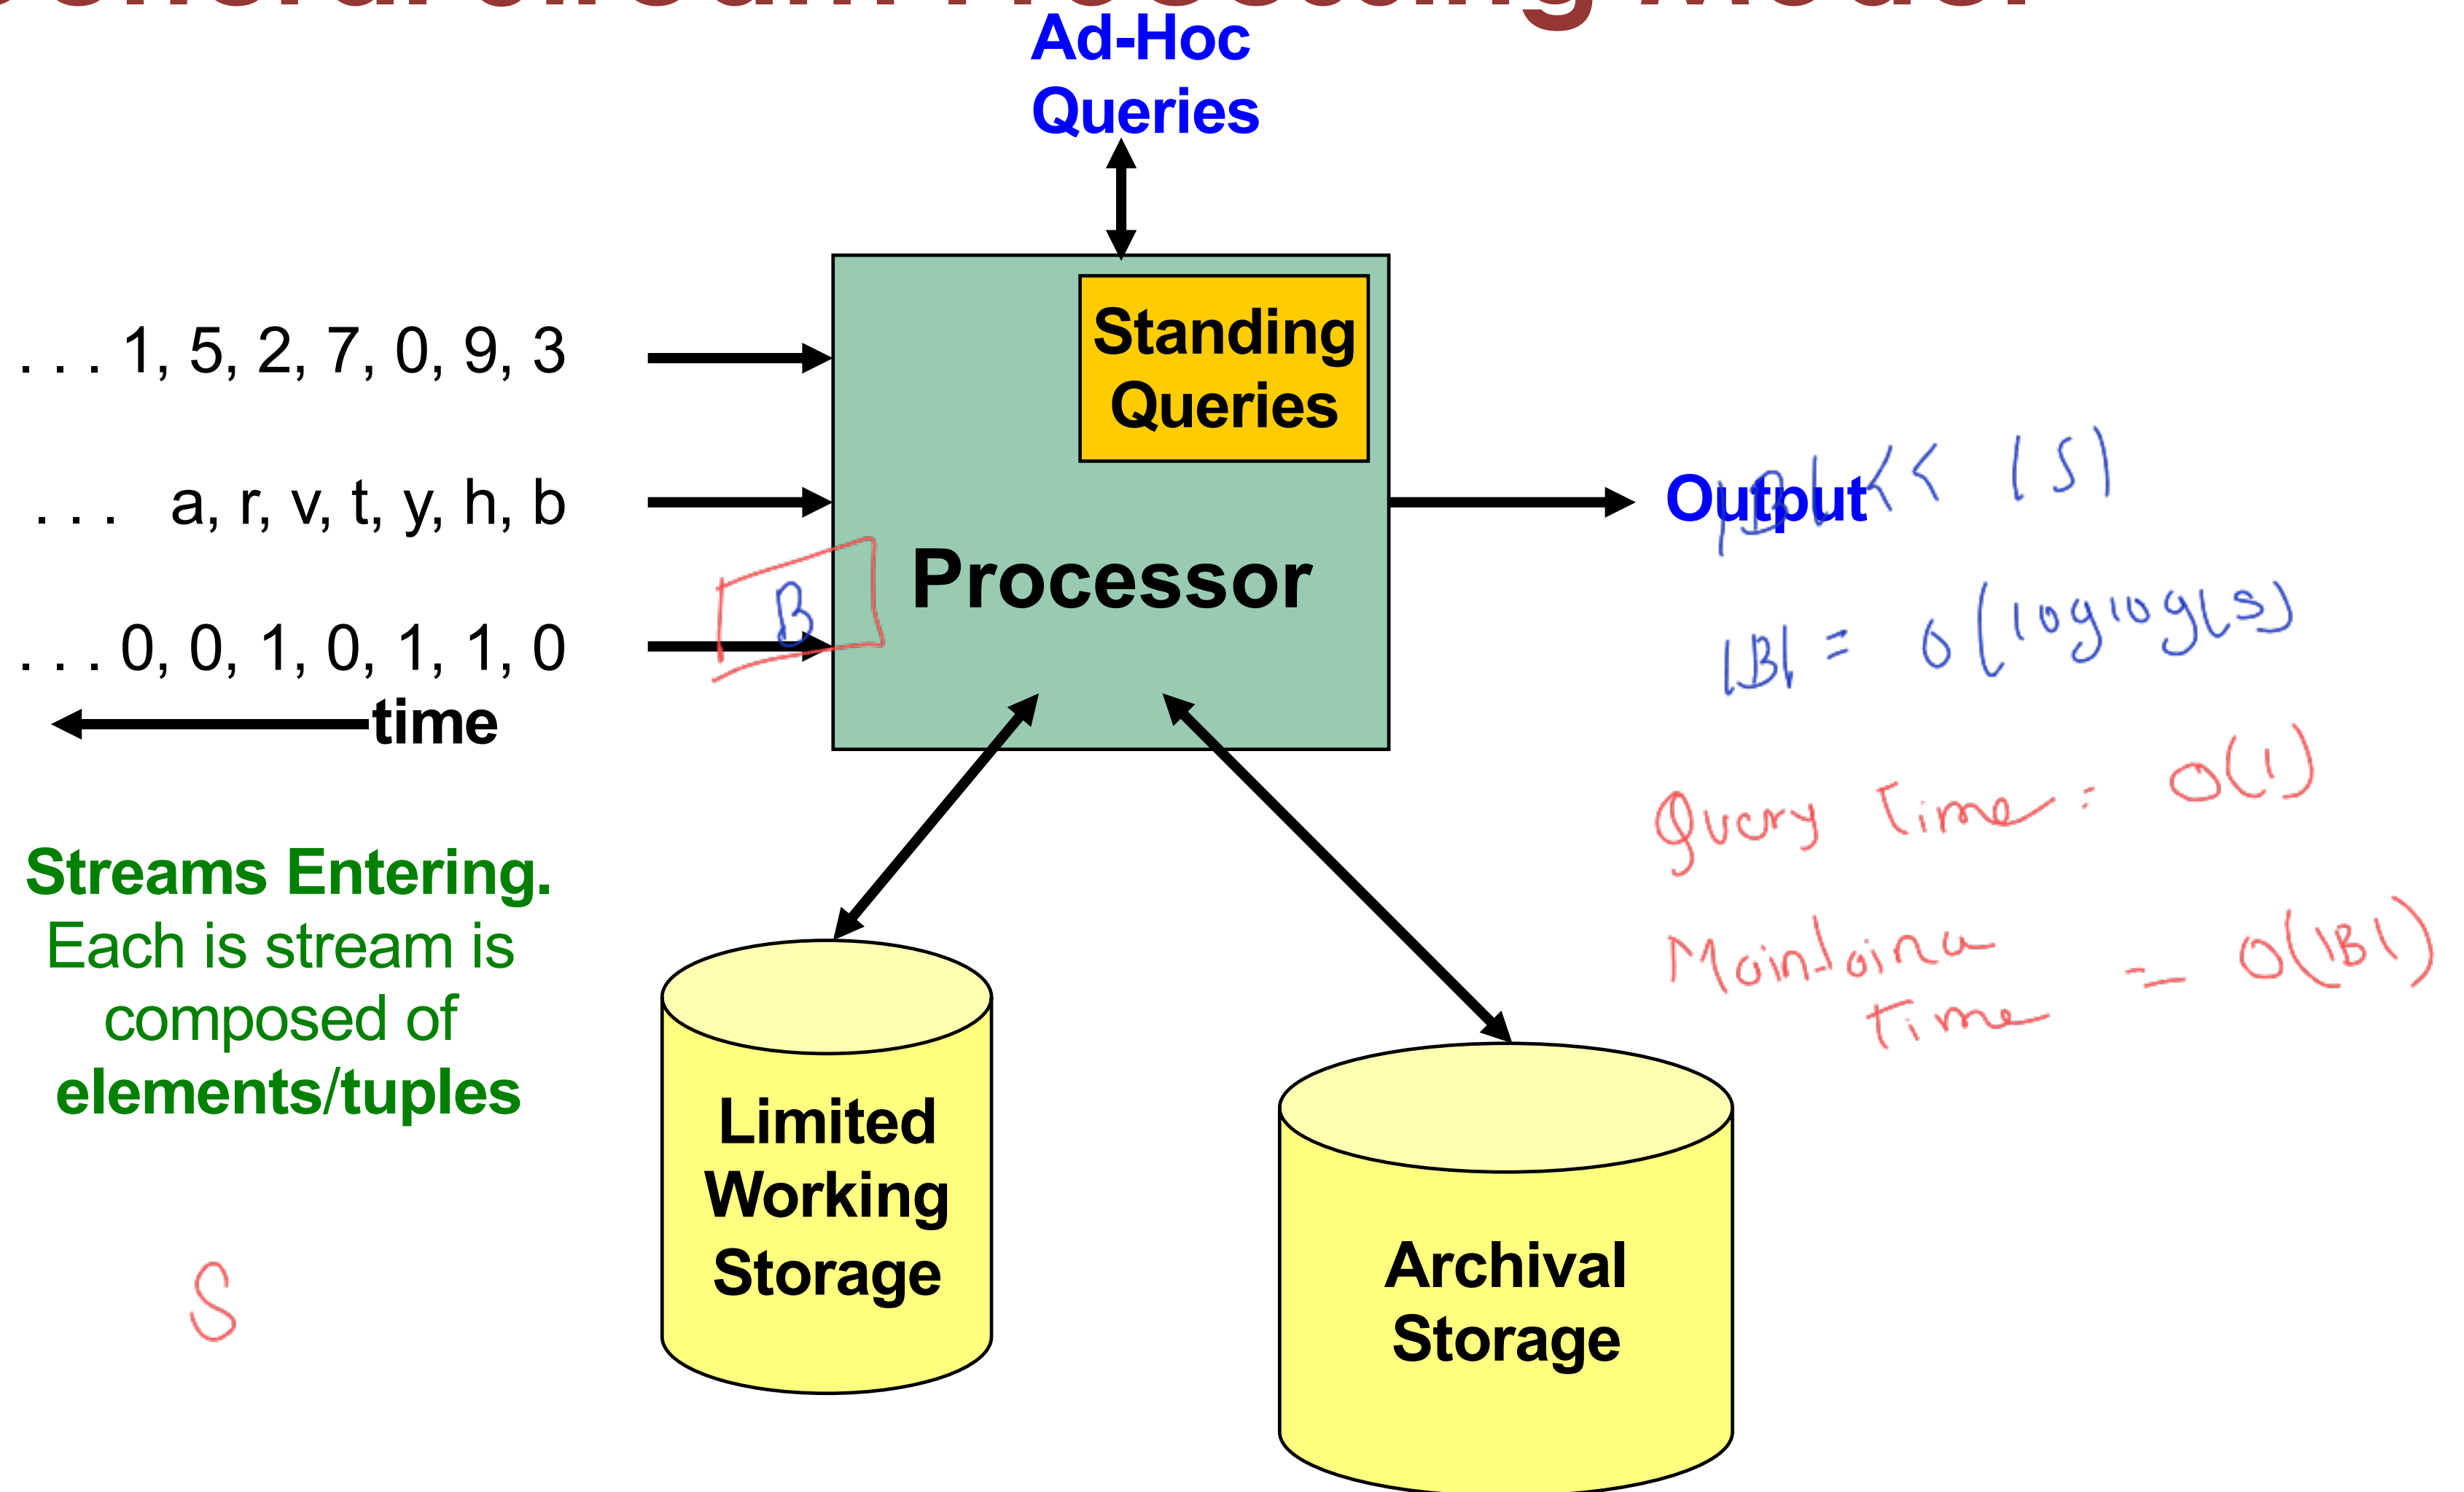
Sourangshu Bhattacharya

Data Streams

- In many data mining situations, we do not know the entire data set in advance
- **Stream Management** is important when the input rate is controlled externally:
 - Google Trends
 - Twitter or Facebook status updates
- We can think of the **data** as **infinite** and **non-stationary** (the distribution changes over time)

- Input **elements** enter at a rapid rate,
at one or more input ports (i.e., **streams**)
 - **We call elements of the stream tuples**
- **The system cannot store the entire stream
accessibly**
- **Q: How do you make critical calculations about
the stream using a limited amount of
(secondary) memory?**

General Stream Processing Model




Reservoir Sampling

- **Problem: Fixed-size sample**
- **Suppose we need to maintain a random sample S of size exactly s tuples**
 - E.g., main memory size constraint
- **Why?** Don't know length of stream in advance
- **Suppose at time n we have seen n items**
 - Each item is in the sample S with equal prob. s/n

How to think about the problem: say $s = 2$

Stream: a x c y z k c d e g...



At $n = 5$, each of the first 5 tuples is included in the sample S with equal prob.

At $n = 7$, each of the first 7 tuples is included in the sample S with equal prob.

Impractical solution would be to store all the n tuples seen so far and out of them pick s at random

Solution: Fixed Size Sample

- **Algorithm (a.k.a. Reservoir Sampling)**
 - Store all the first s elements of the stream to S
 - Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s/n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random
- **Claim:** This algorithm maintains a sample S with the desired property:
 - After n elements, the sample contains each element s so far with probability s/n

Proof: By Induction

- **We prove this by induction:**
 - Assume that after n elements, the sample contains each element seen so far with probability s/n
 - We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element seen so far with probability s/n
- **Base case:**
 - After we see $n=s$ elements the sample S has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction

- **Inductive hypothesis:** After n elements, the sample S contains each element seen so far with prob. s/n
- **Now element $n+1$ arrives**
- **Inductive step:** For elements already in S , probability that the algorithm keeps it in S is:

$$\left(1 - \frac{s}{n+1}\right) + \left(\frac{s}{n+1}\right)\left(\frac{s-1}{s}\right) = \frac{n}{n+1}$$

Element $n+1$ discarded

Element $n+1$
not discarded

Element in the
sample not picked

- So, at time n , tuples in S were there with prob. s/n
- Time $n \rightarrow n+1$, tuple stayed in S with prob. $n/(n+1)$
- So prob. tuple is in S at time $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

Type of sketch

- Sample → Reservoir Sampling.
- Hash Table
 - ↳ False positive
- Specialized data structure

Queries

- Membership
- Count distinct
- min/max/avg/median
- moment.

Criteria should be satisfied

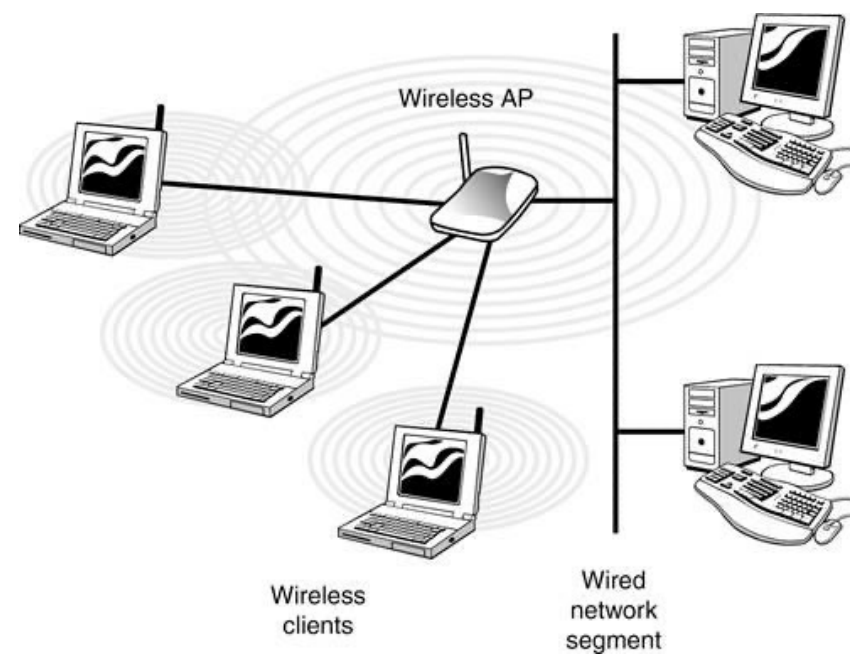
$$P(|g(S) - g(B)| < \epsilon) > 1 - \delta$$

$$|B| \ll |S|$$

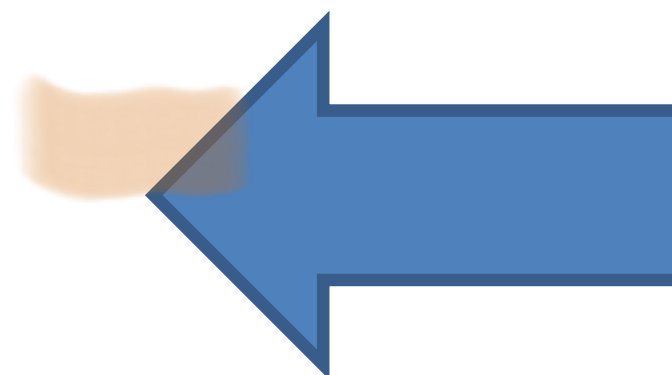
Bloom Filters



ISBN present in collection? *given dataset Z, whether Z appear in S (streaming dataset)*



IP seen by switch?



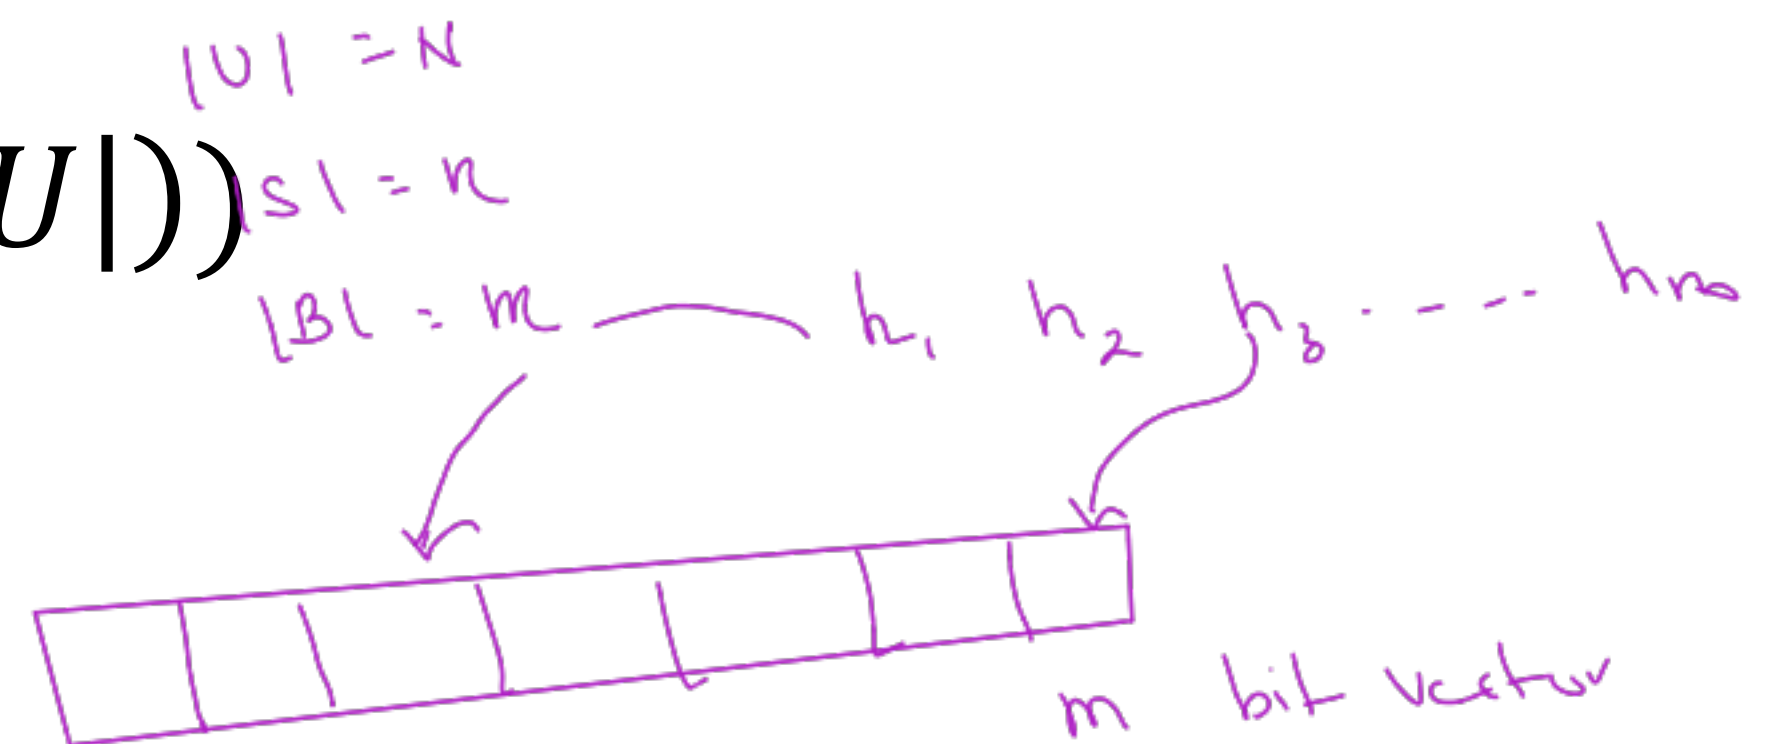
10.0.21.102

- Universe U , but need to store a set of n items, $n \ll |U|$.

- Hash table of size m :

- Space $O(m + n \log(|U|))$

★ – Query time $O(\frac{n}{m})$



if z has occurred in S
 $q(z)$ is always 1

Sometimes if z has not occurred
 still $q(z)$ can be 1 (False positive)

- Universe U , but need to store a set of n items, $n \ll |U|$.
- Hash table of size m :
 - Space $O(m + n \log(|U|))$
 - Query time $O(\frac{n}{m})$
- Bit array of size $|U|$
 - Space $|U|$.
 - Query time $O(1)$.

$P(\text{False Positive}) = ?$
 $P(\text{False Negative}) = 0$

* False positive happens because of over-writing to 1
 * Find the probability when one of hash function overwrites 1 while inserting x_{n+1}

* What is probability that $B[h_1(x_{n+1})] = 1$ already.

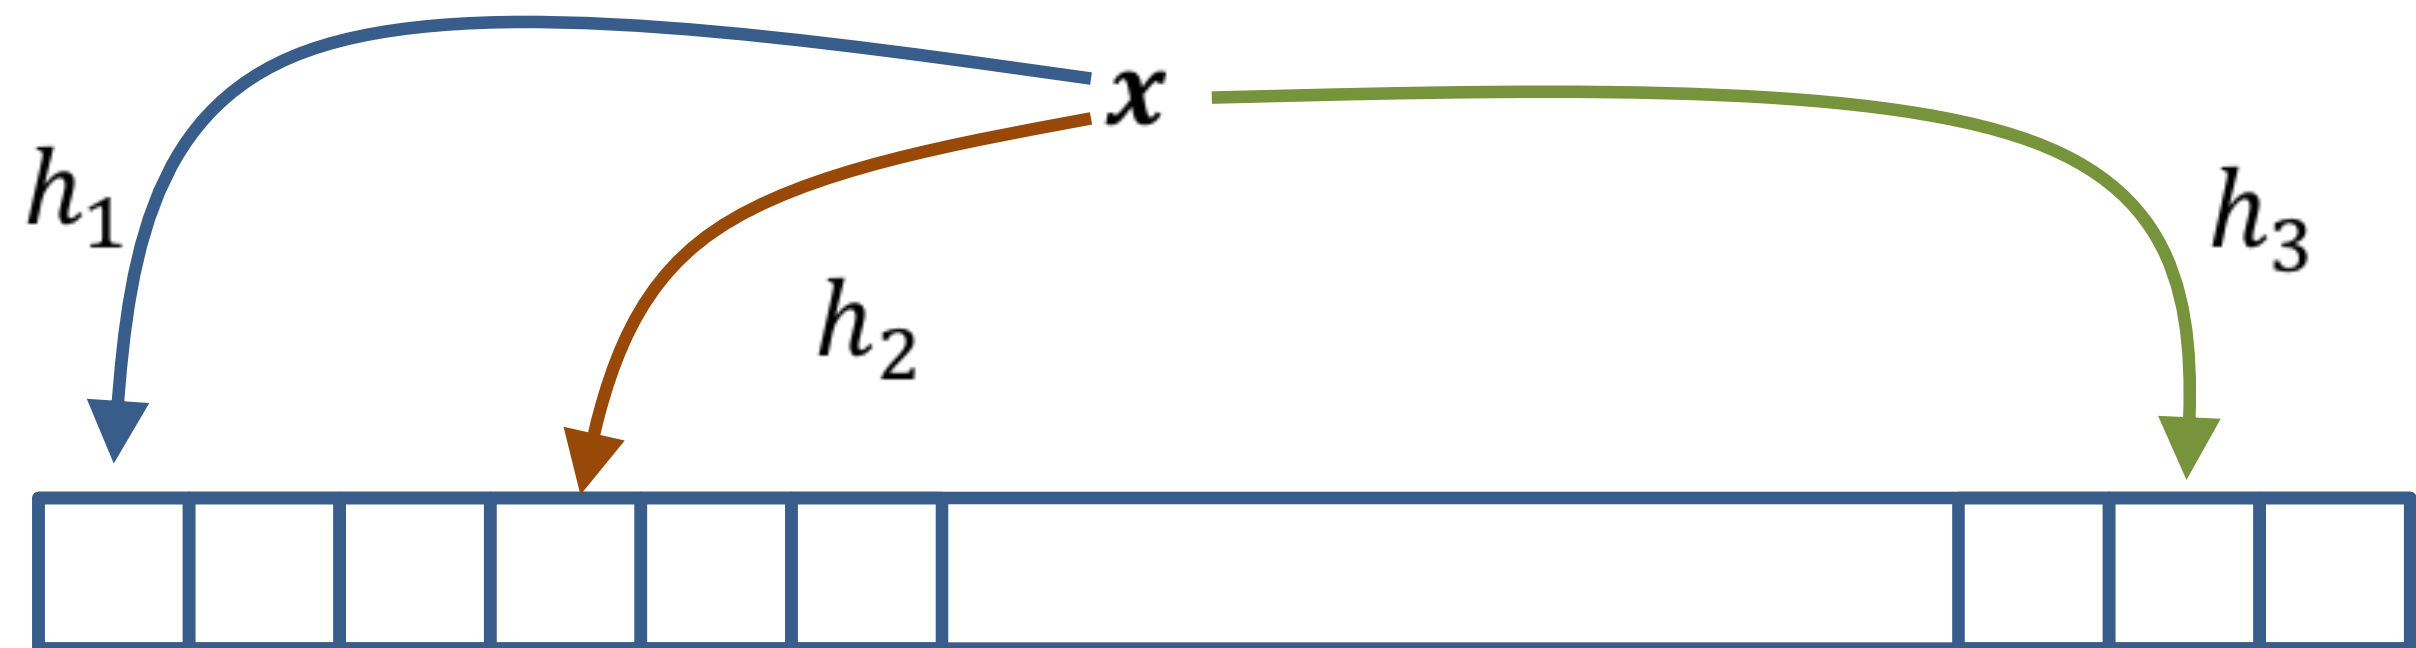
* Find the expected number of #1's in B after n -operation.

$$= \frac{\# \text{1's } B}{|B|}$$

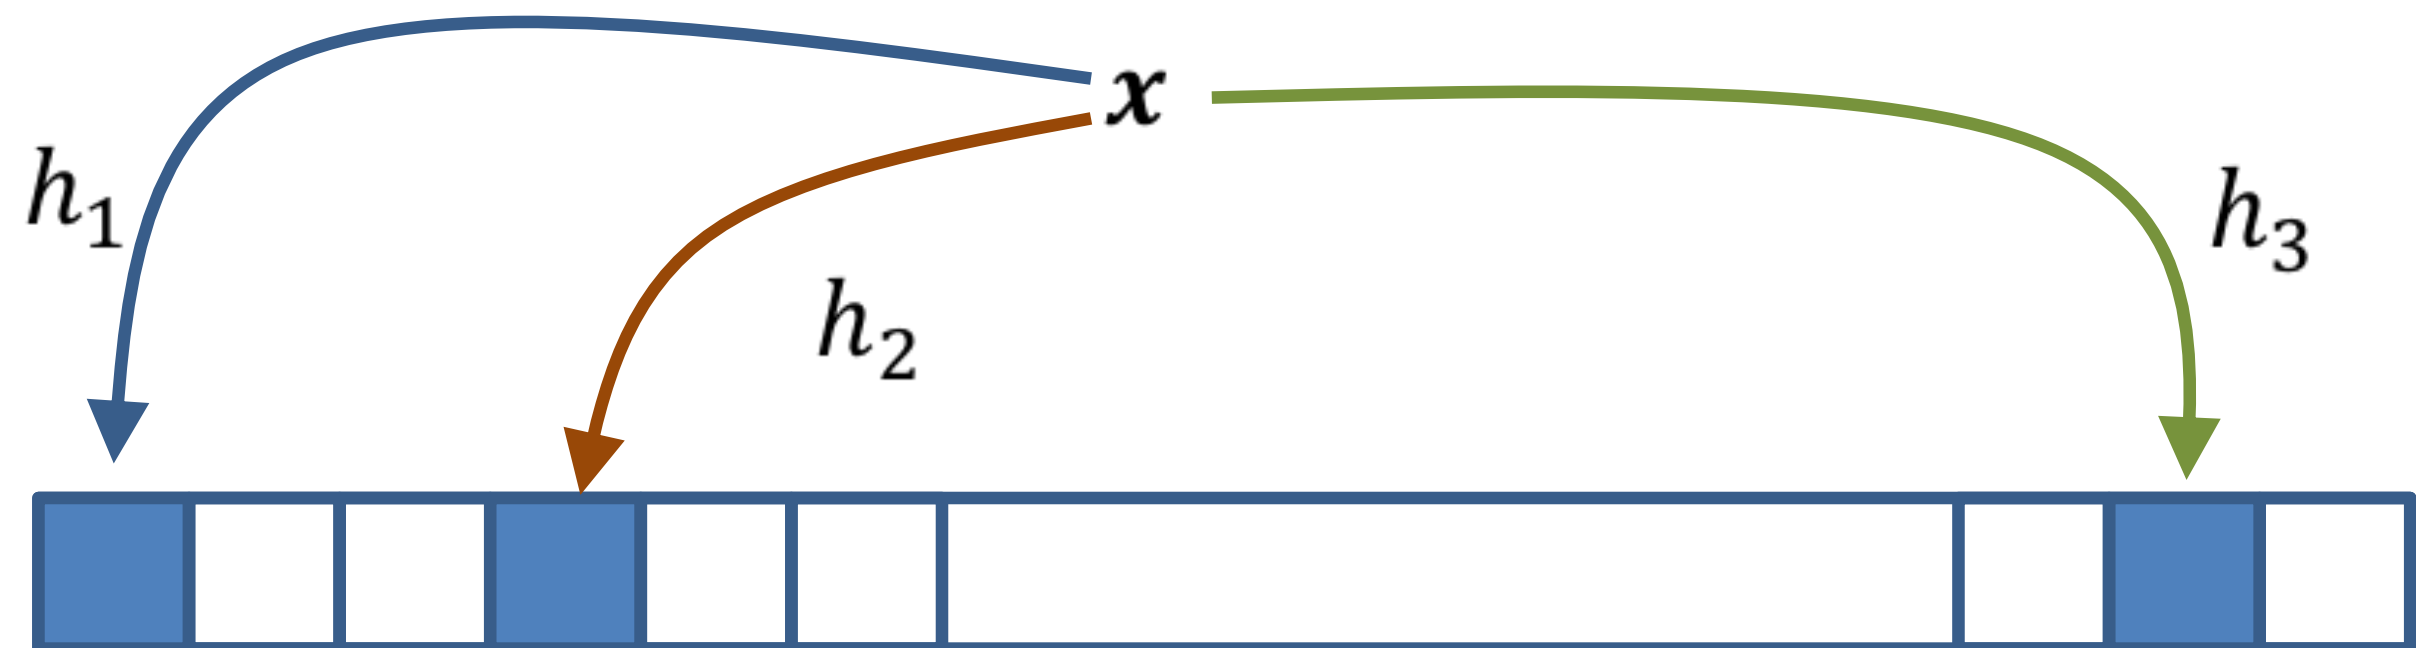
- In hash table construction, we used random hash functions
 - we never return incorrect answer
 - query time is a random variable
 - These are Las Vegas algorithms
- In Monte-Carlo randomized algorithms, we are allowed to return incorrect answers with (small) probability, say, δ

[Bloom, 1970]

- A bit-array B , $|B| = m$
- k hash functions, h_1, h_2, \dots, h_k , each $h_i \in U \rightarrow [m]$



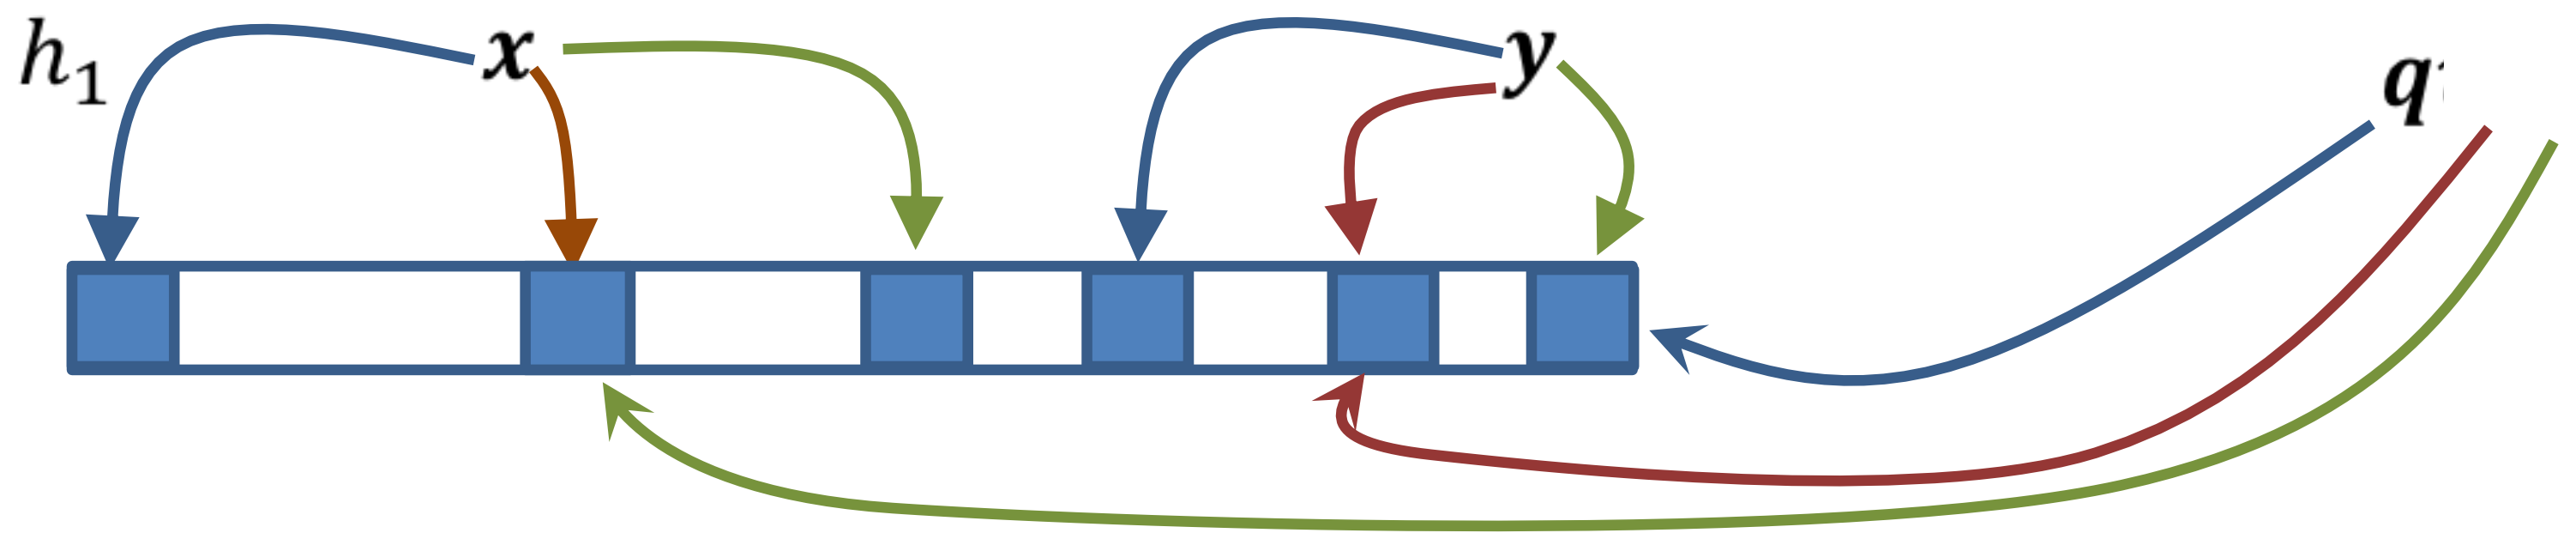
- A bit-array B , $|B| = m$
- k hash functions, h_1, h_2, \dots, h_k , each $h_i \in U \rightarrow [m]$



- *Initialize*(B)
 - for $i \in \{1, \dots, m\}$, $B[i] = 0$
- *Insert* (B, x)
 - for $i \in \{1, \dots, k\}$, $B[h_i(x)] = 1$
- *Lookup* (B, x)
 - If $\bigwedge_{i \in \{1, \dots, k\}} B[h_i(x)]$, return PRESENT, else ABSENT

- If the element x has been added to the Bloom filter, then $Lookup(B, x)$ always return **PRESENT**

- If the element x has been added to the Bloom filter, then $Lookup(B, x)$ always return **PRESENT**
- If x has not been added to the filter before?
 - $Lookup$ sometimes still return **PRESENT**



- Want to minimize the probability that we return false positive
- Parameters $m = |B|$ and $k =$ number of hash functions
- $k = 1 \Rightarrow$ normal bit-array
- What is effect of changing k ?

- Increasing k
 - Possibly makes it harder for false positives to happen in *Lookup* because of $\bigwedge_{i \in \{1, \dots, k\}} B[h_i]$
 - But also increases the number of filled up positions
- We can analyse to find out an “optimal k ”

- $m = |B|$, n elements inserted
- If x has not been inserted, what is the probability that $Lookup(B, x)$ returns PRESENT?

- $m = |B|$, n elements inserted
- If x has not been inserted, what is the probability that $Lookup(B, x)$ returns PRESENT?
- Assume $\{h_1, h_2, \dots, h_k\}$ are independent and $\Pr[h_i(\cdot) = j] = \frac{1}{m}$ for all positions j

↳ equal probability for all m blocks

- Probability of a bit being zero:

$$P[B_j = 0] = \left(1 - \frac{1}{m}\right)^{kn} \approx e^{-\frac{kn}{m}}$$

- The expected number of zero bits is given by:
 $me^{-kn/m}$.

- $P[\text{lookup}(B, x) = \text{PRESENT}] = \left(1 - e^{-\frac{kn}{m}}\right)^k$
Probability that
Particular B_j
for h_j
already set.

- We can choose k to minimize this probability.
Probability that x is already inserted into B

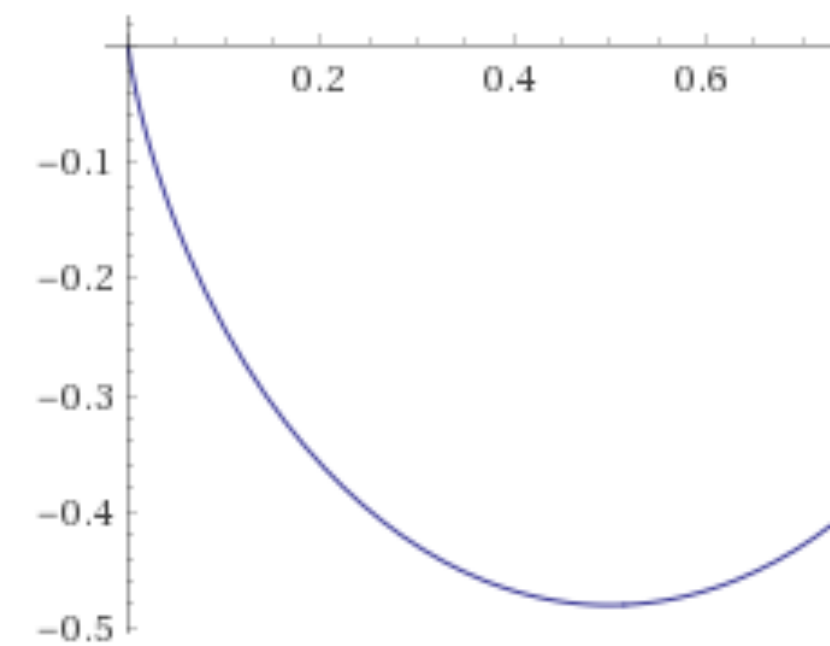
functions

- $p = e^{-kn/m}$

- Log (False Positive) =

$$\log(1 - p)^k = k \log(1 - p) = -\frac{m}{n} \log(p) \log(1 - p)$$

Minimized at $p = \frac{1}{2}$, i.e. $k = m \log(2)/n$



Computed by Wolfram|Alpha

- This “optimal” choice gives false positive = $2^{-m} \log(2)/n$
- If we want a false positive rate of δ , set $m = \left\lceil \frac{\log\left(\frac{1}{\delta}\right)n}{\log^2(2)} \right\rceil$

Example: If we want 1% FPR, we need 7 hash functions and total $10n$ bits

- Widespread applications whenever small false positives are tolerable
- Used by browsers
 - to decide whether an URL is potentially malicious: a BF is used in browser, a positives are actually checked with the server.
- Databases e.g. BigTable, HBase, Cassandra, Postgrepsql use to avoid disk lookups for non-existent rows/columns
- Bitcoin for wallet synchronization....

- Chief drawback is that BF does not allow deletions

[Fan et al 00]

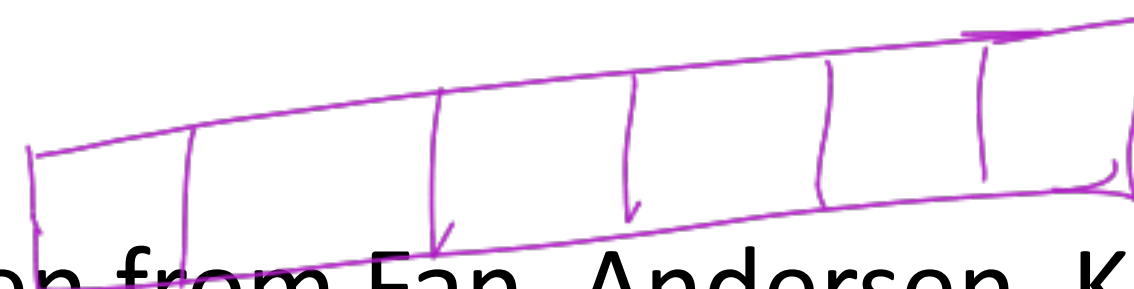
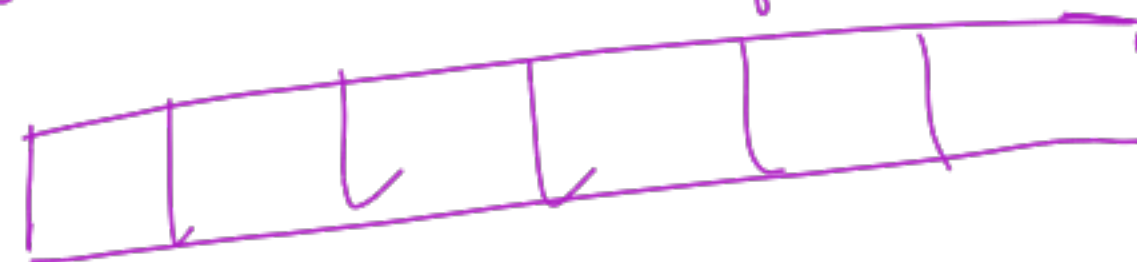
- Counting Bloom Filter
 - Every entry in BF is a small counter rather than a single bit
 - *Insert*(x) increments all counters for $\{h_i(x)\}$ by 1
 - *Delete*(x) decrements all $\{h_i(x)\}$ by 1
 - maintains 4 bits per counter
 - False negatives can happen, but only with low probability

CUCKOO FILTERS

So if for x_{n+1} $h_1(x_{n+1})$ is already
1. Go for $g_1(h(x_{n+1}))$ if that is not
1 - Set it to 1. if one does something

maintains two filters

h_1, \dots, h_k



g_1, \dots, g_k

Slides taken from Fan, Andersen, Kaminsky, M

- Fingerprint(x): A hash value of x
 - Lower false positive rate ϵ , longer fingerprint

→ Problem: on given n -data points from N Count the number of distinct numbers
Can do in $O(\log N)$?

given data point

$x \longrightarrow i \quad (\log_2 N \text{ bits})$

$x_i \longrightarrow h(x_i) \longrightarrow \text{binary} \longrightarrow 11000_3$
 i (trailing 0s in x)

Calculate for all i and

if \max is $\max - i$

so the number of distinct will be $2^{\max - i}$

$R \leq \max - i$

Flojole-Martin
Algorithm

0:
1:
2:
3:
4:
5:
6:
7:

$$m(1-\epsilon)^n < 2^R < m(1+\epsilon)^n$$

R : max number of
trillion zeros in
any x

- **Fingerprint(x)**: A hash value of x
 - Lower false positive rate ϵ , longer fingerprint
- **Insert(x)**:
 - add **Fingerprint(x)** to hash table

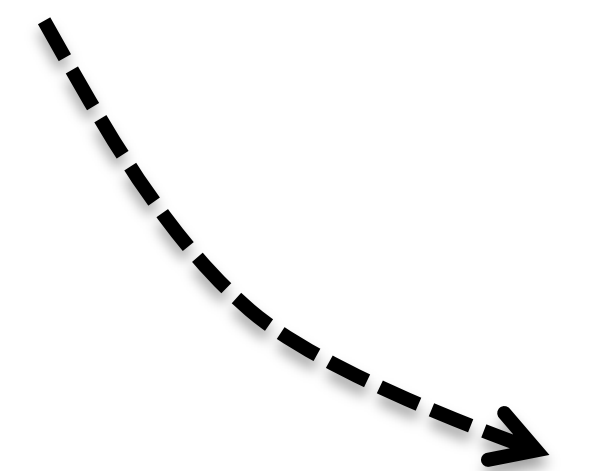
0:
1:
2:
3:
4:
5:
6:
7:

- **Fingerprint(x):** A hash value of x
 - Lower false positive rate ϵ , longer fingerprint

- **Insert(x):**
 - add **Fingerprint(x)** to hash table

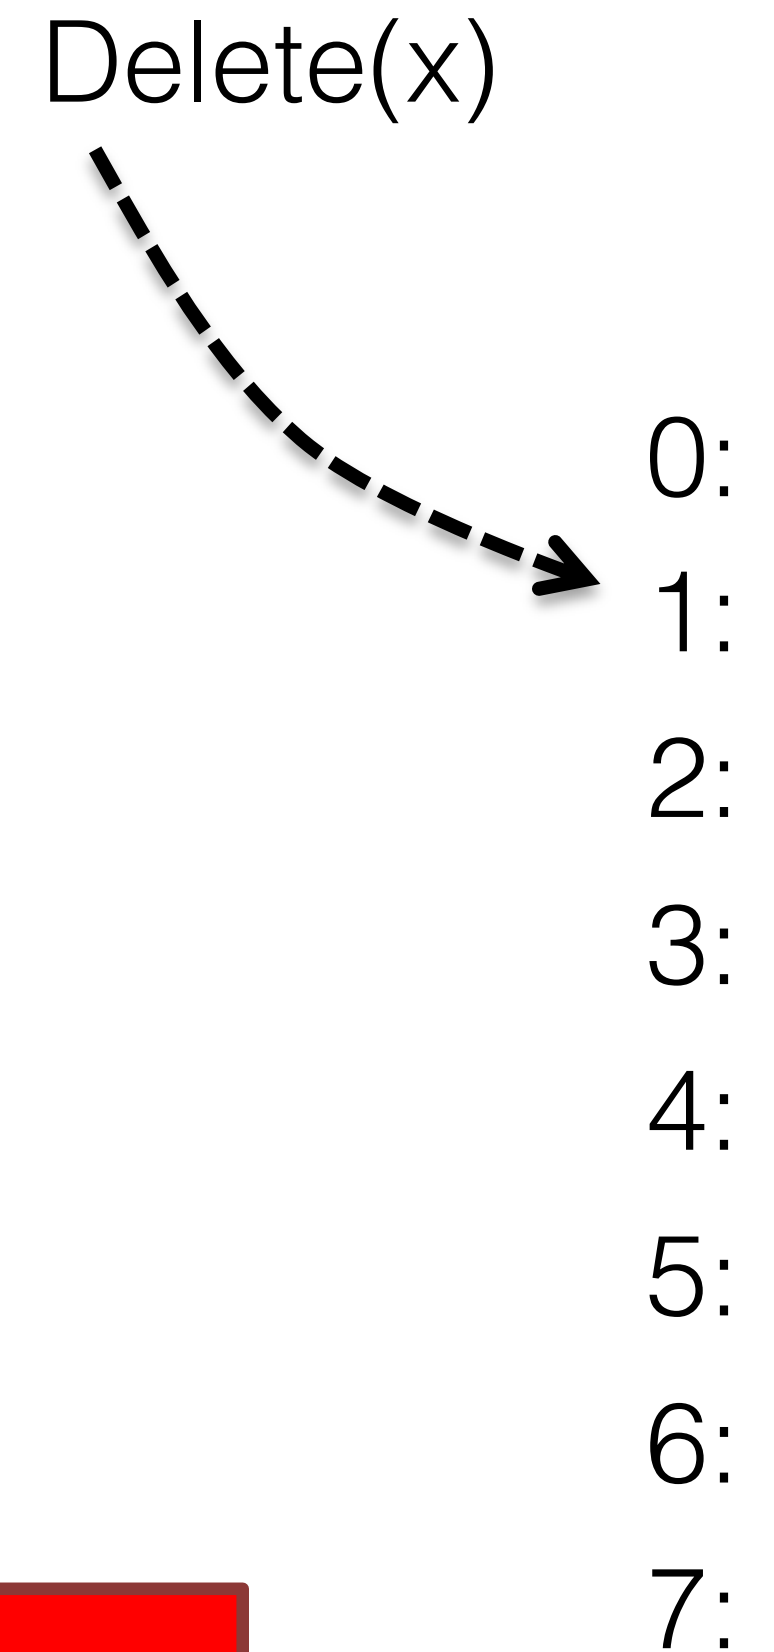
- **Lookup(x):**
 - search **Fingerprint(x)** in hashtable

Lookup(x) = four



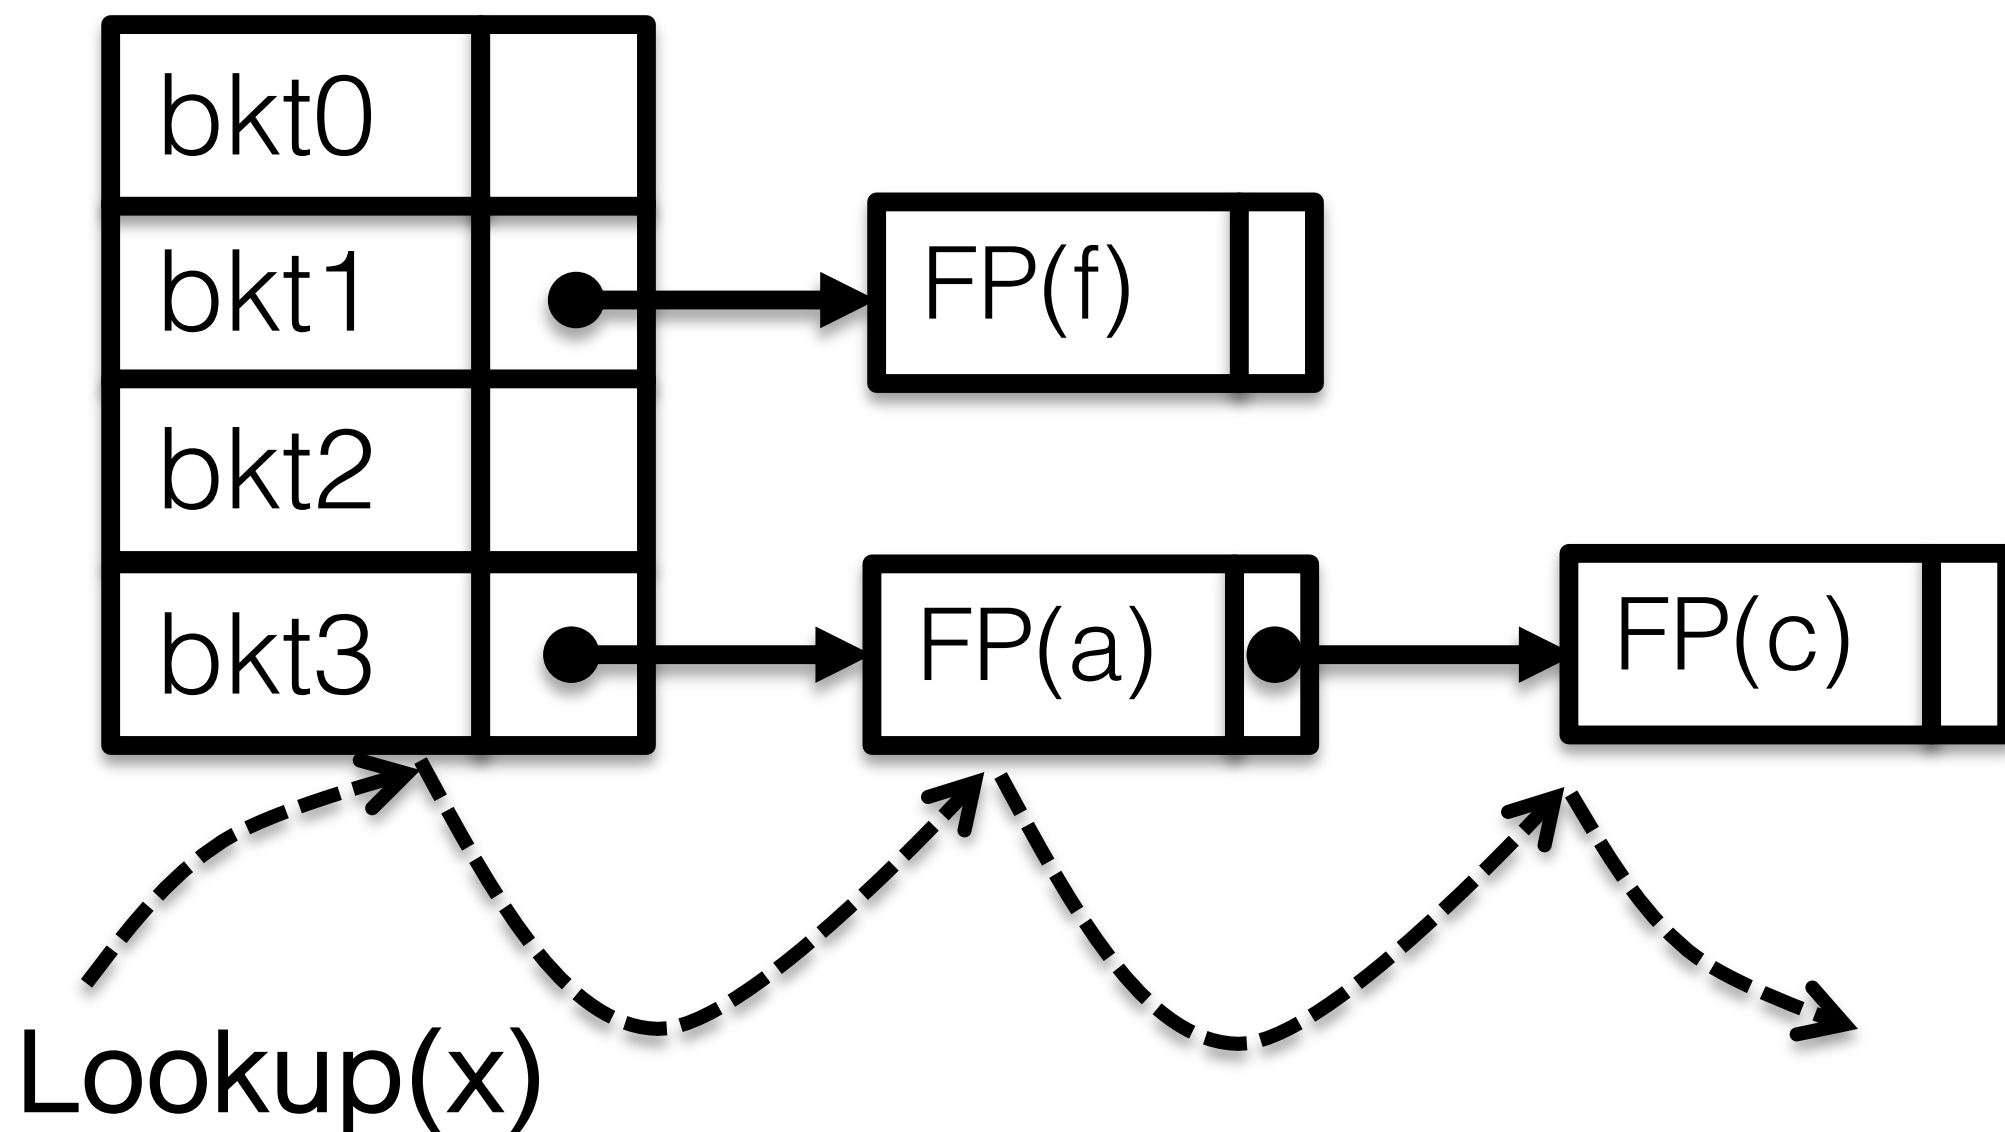
0:
1:
2:
3:
4:
5:
6:
7:

- **Fingerprint(x):** A hash value of x
 - Lower false positive rate ϵ , longer fingerprint
- **Insert(x):**
 - add **Fingerprint(x)** to hash table
- **Lookup(x):**
 - search **Fingerprint(x)** in hashtable
- **Delete(x):**
 - remove **Fingerprint(x)** from hashtable



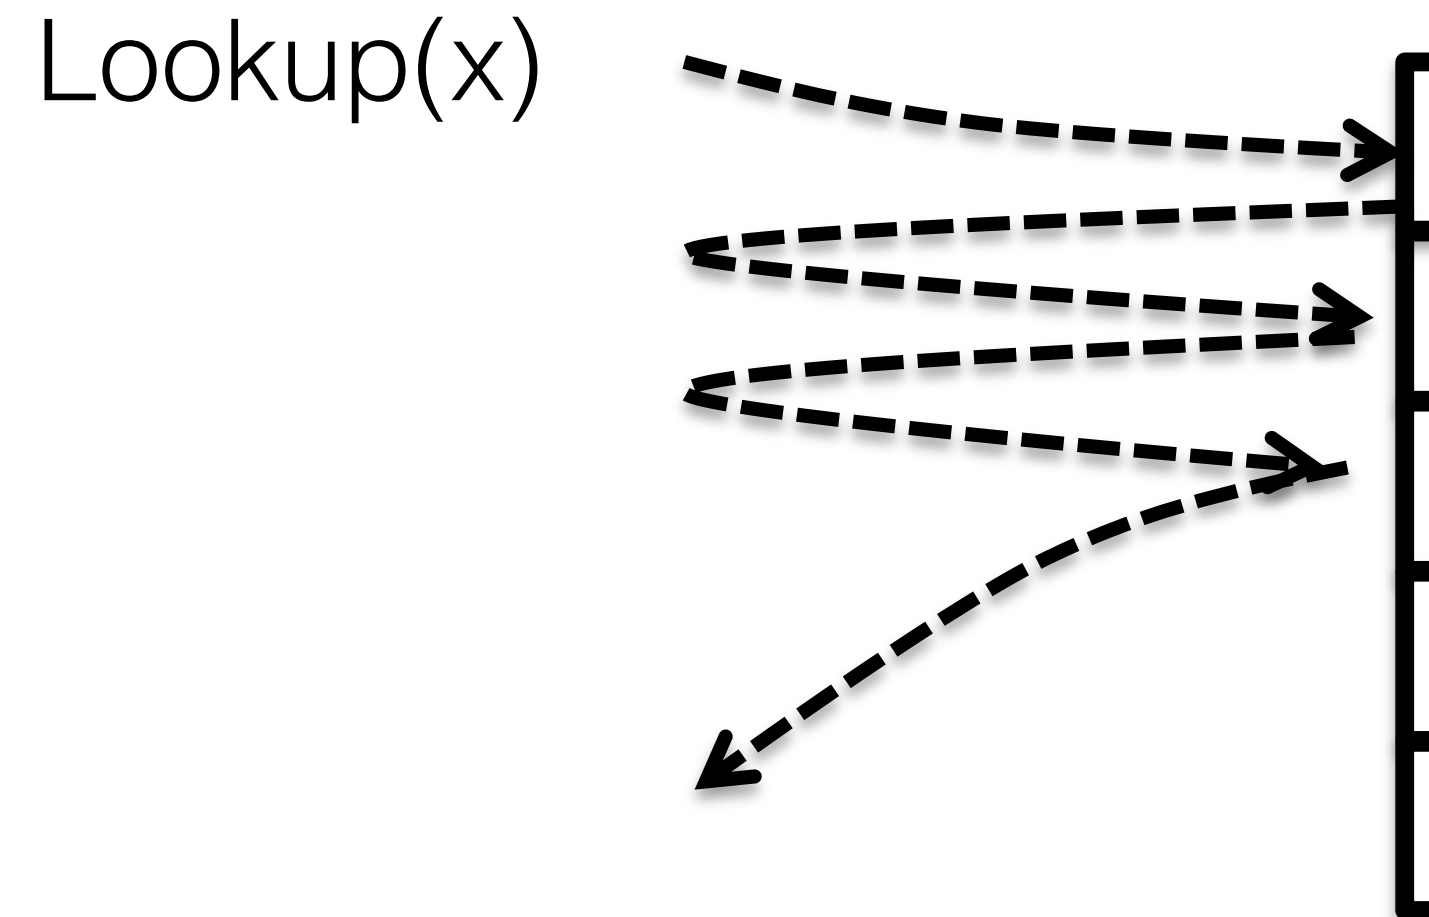
How to Construct Hashtable?

- Chaining :

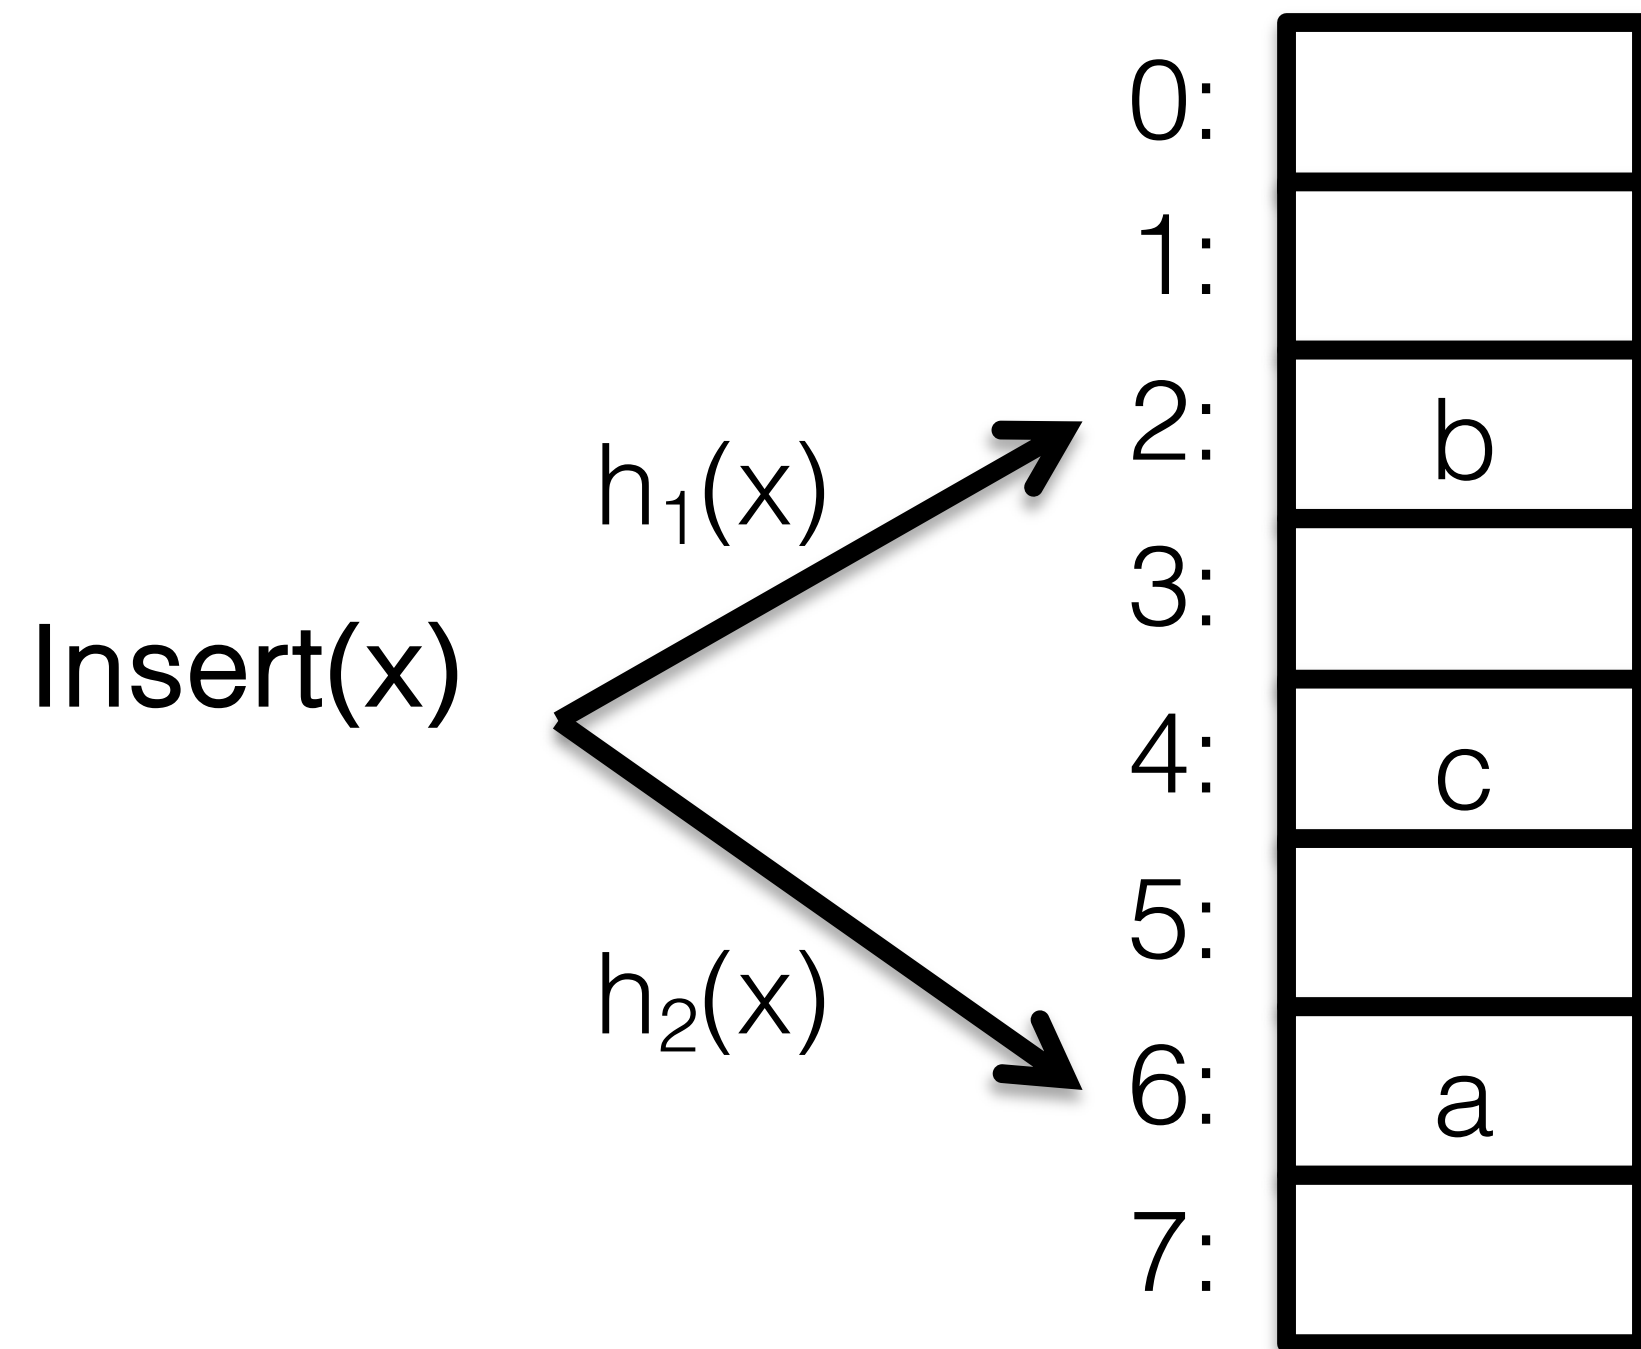


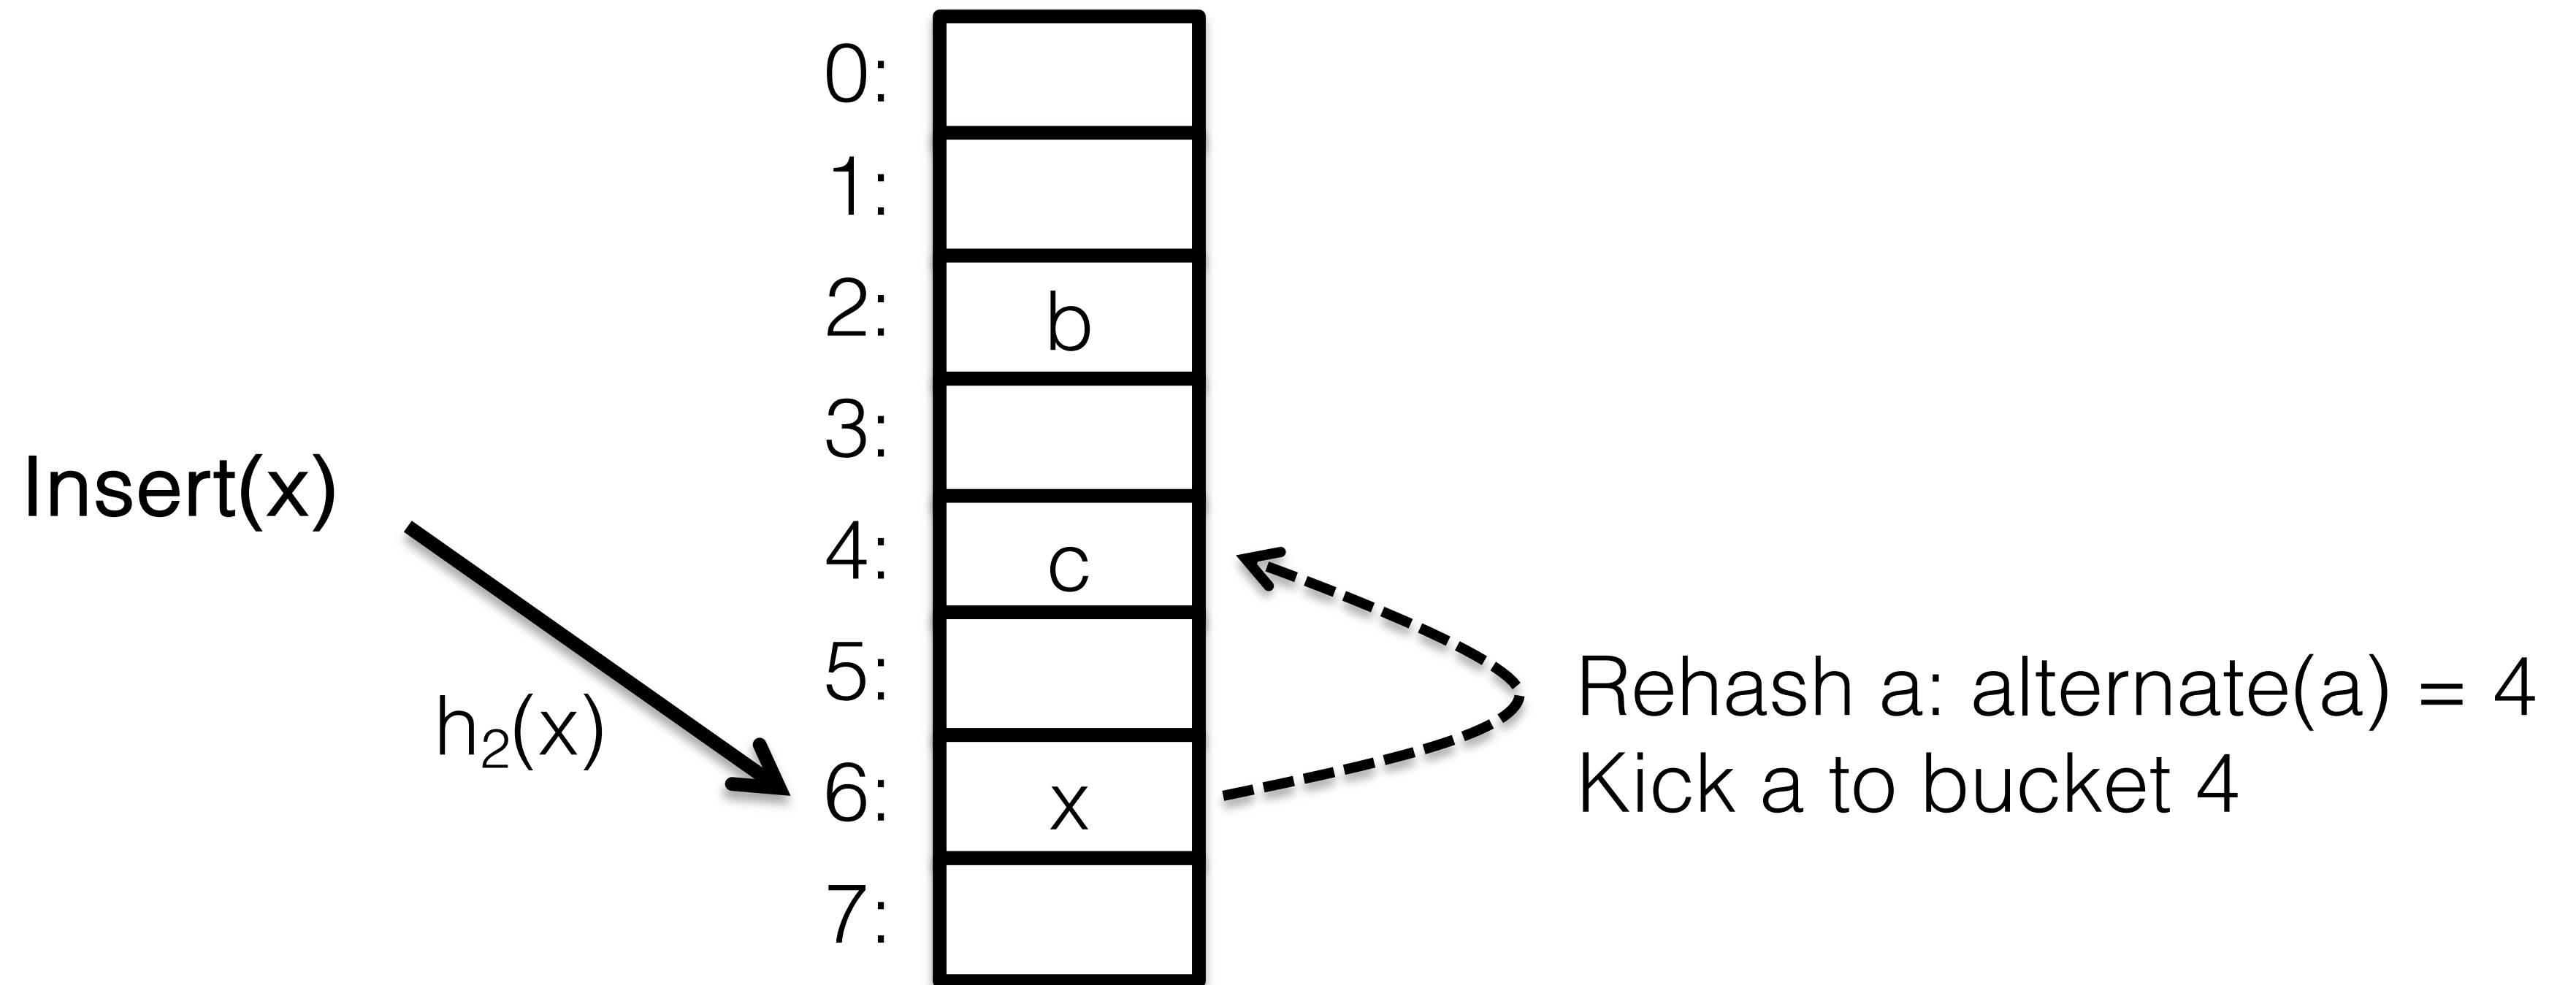
- Pointers →
low space utilization

- Linear Probing

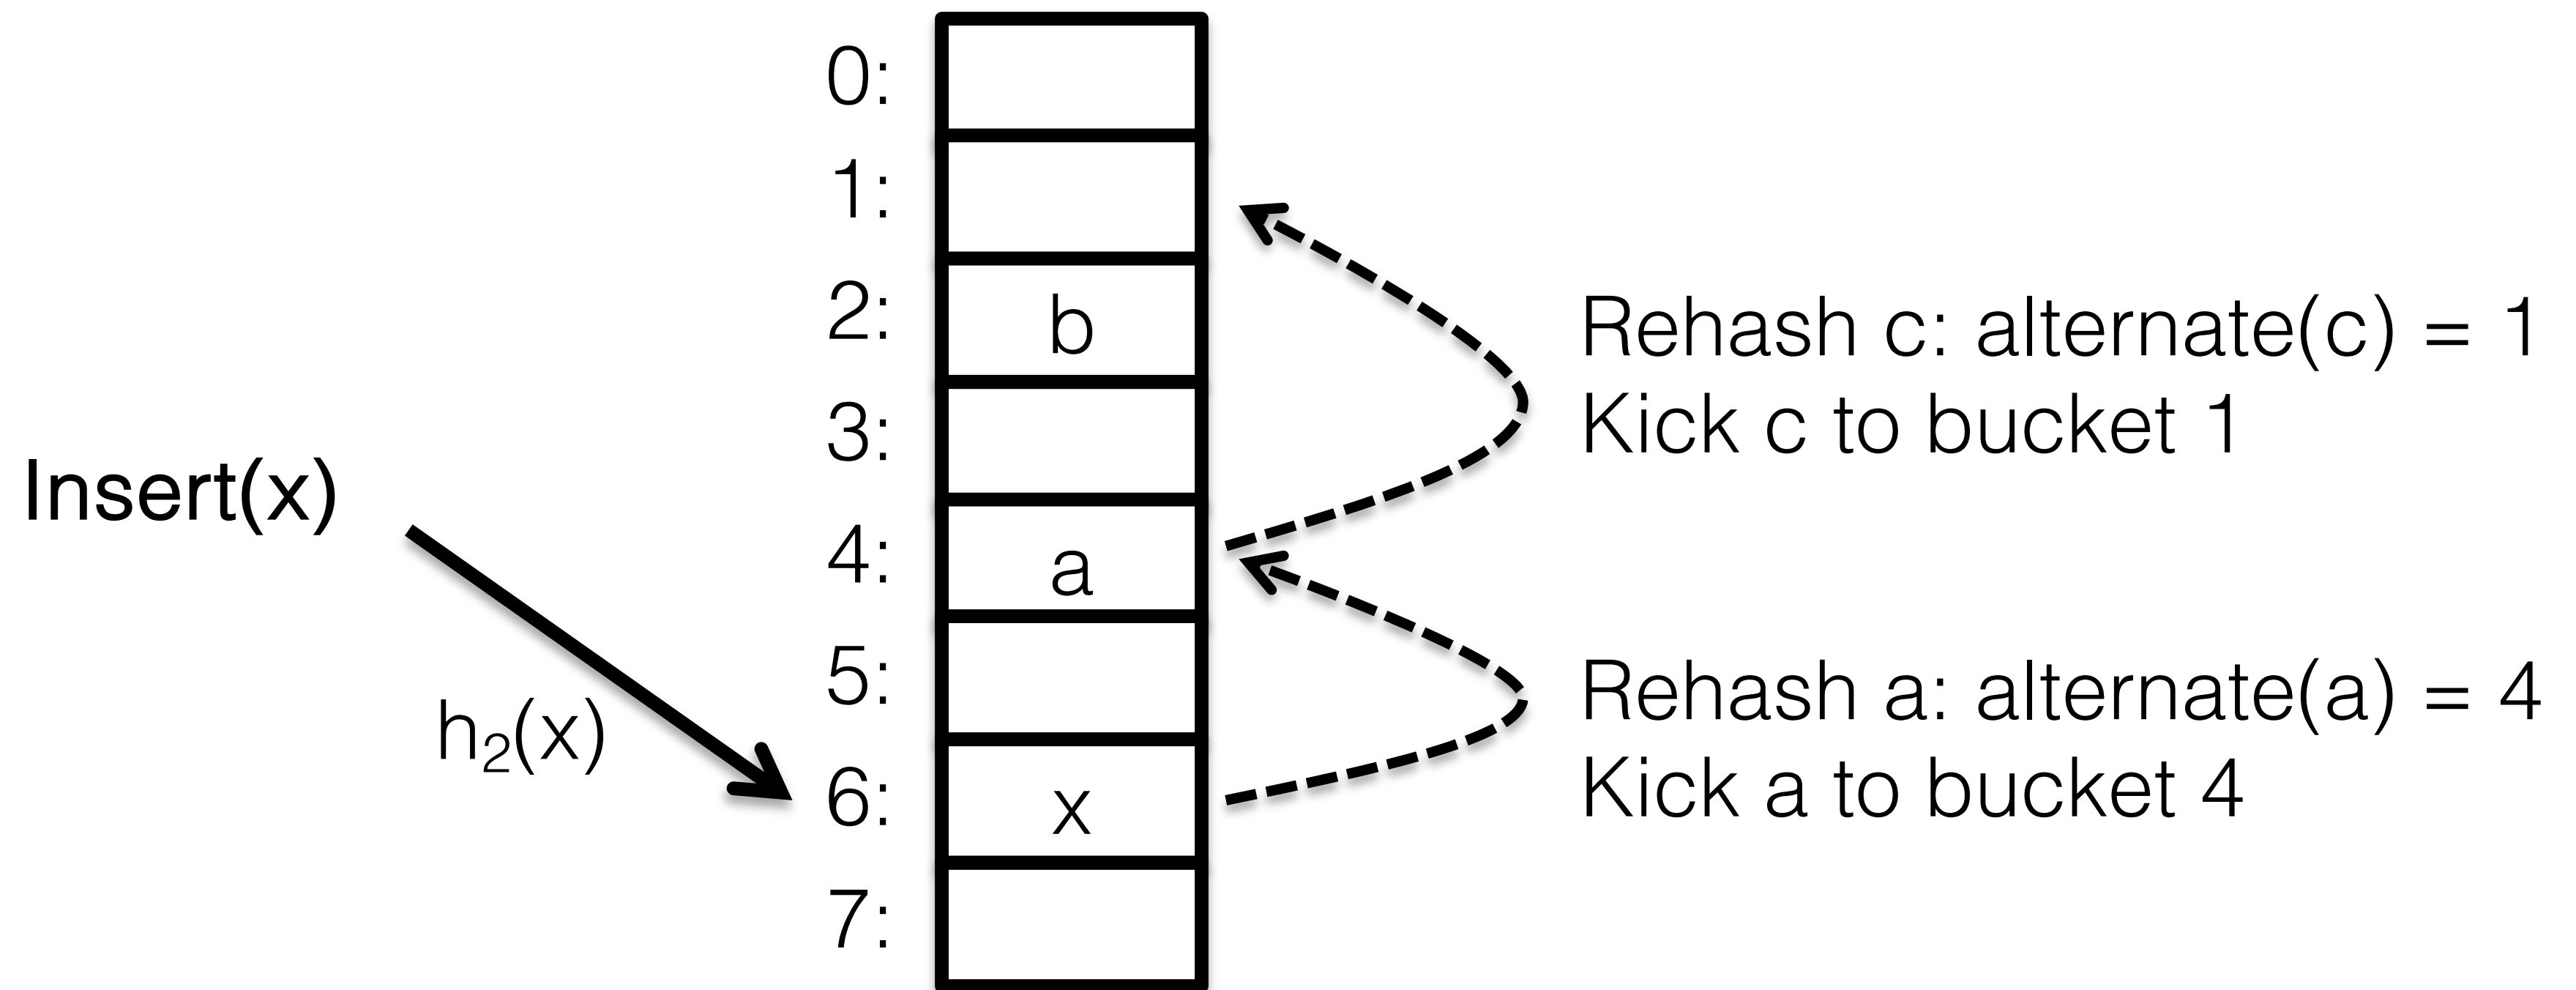


- Making lookups $O(1)$ require
% table empty →
low space utilization
- Compare multiple fingerprints
sequentially →
more false positives

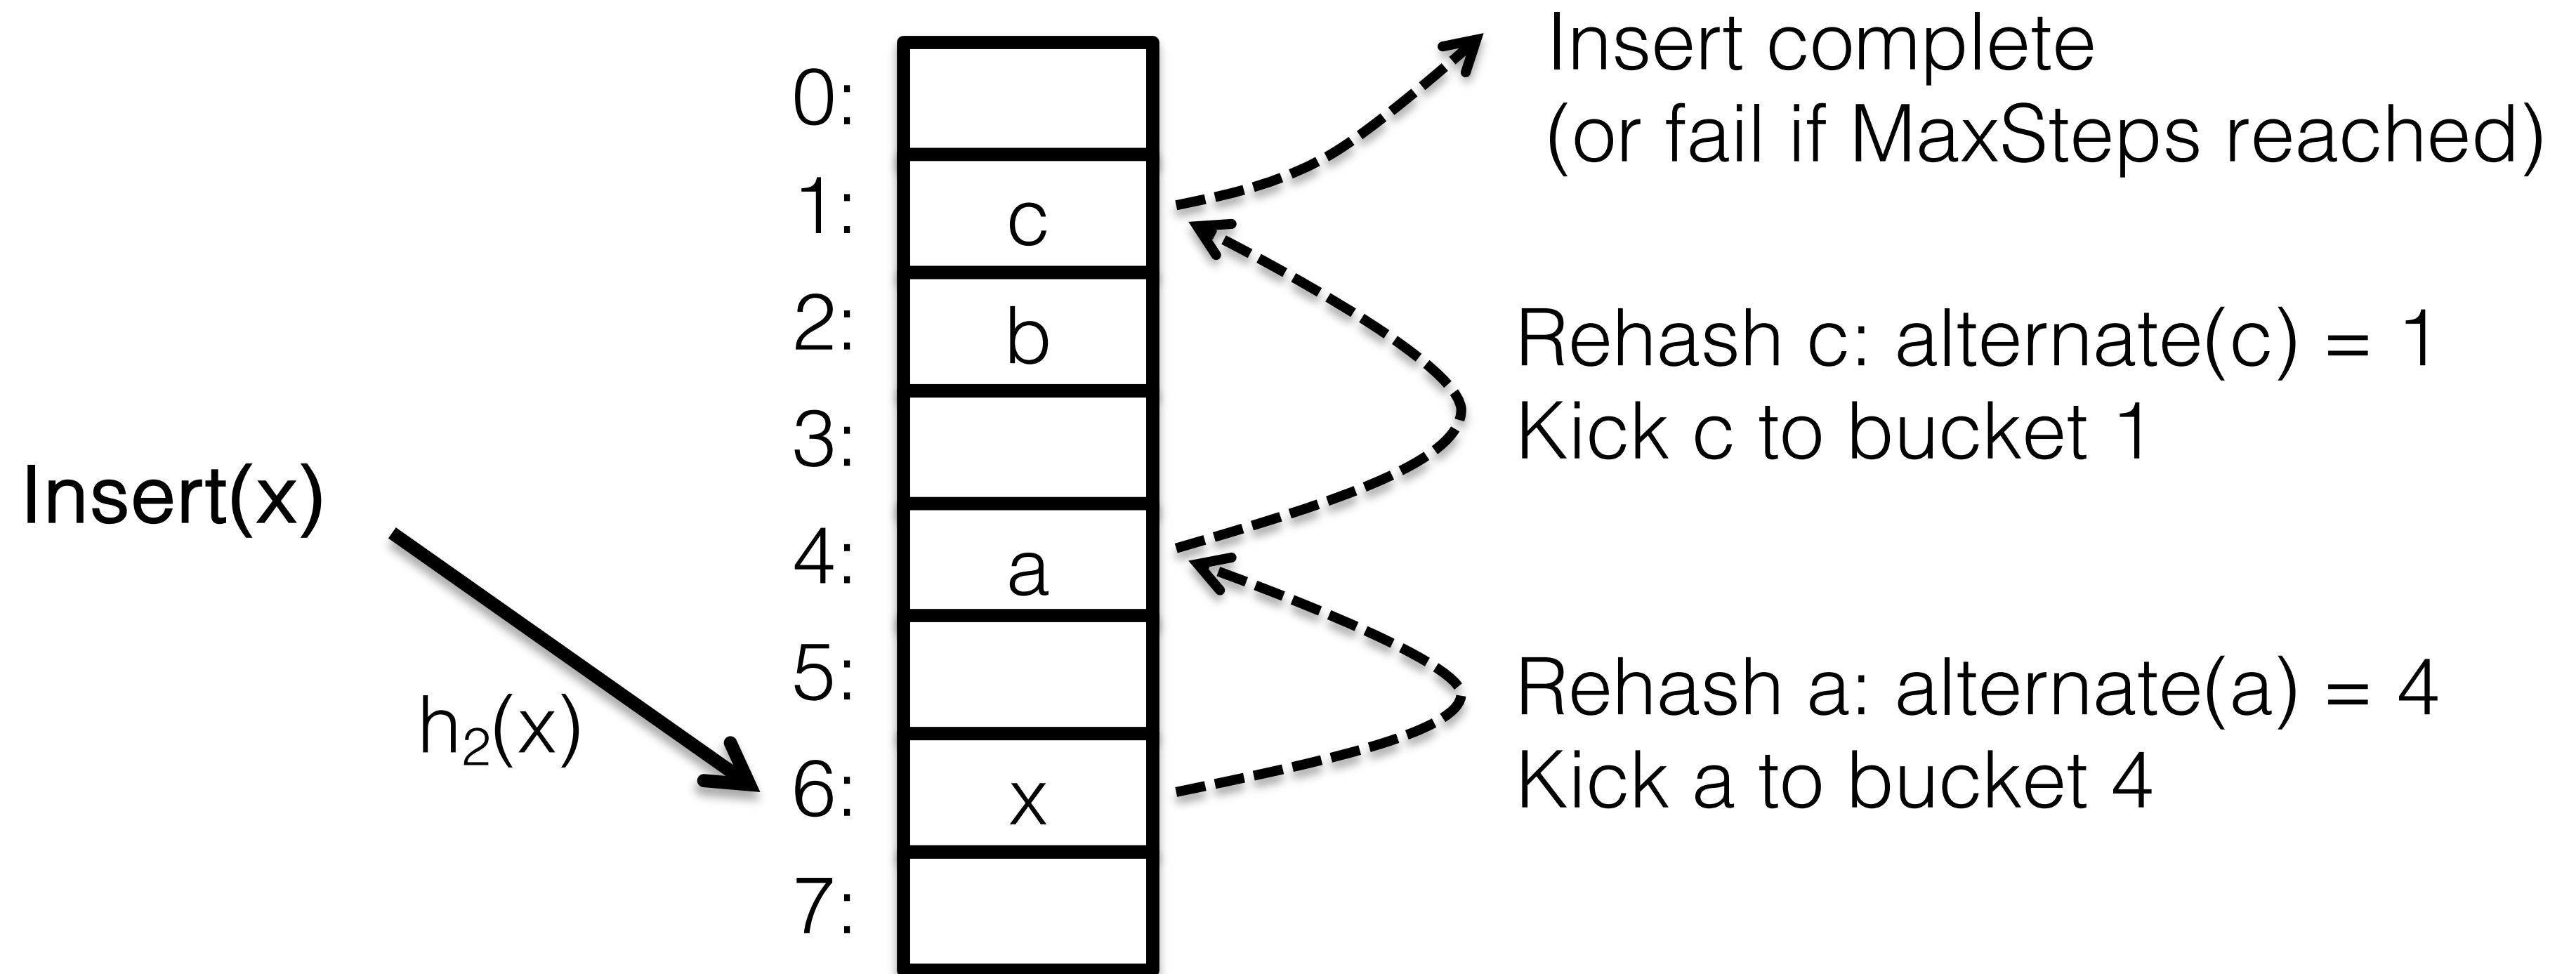




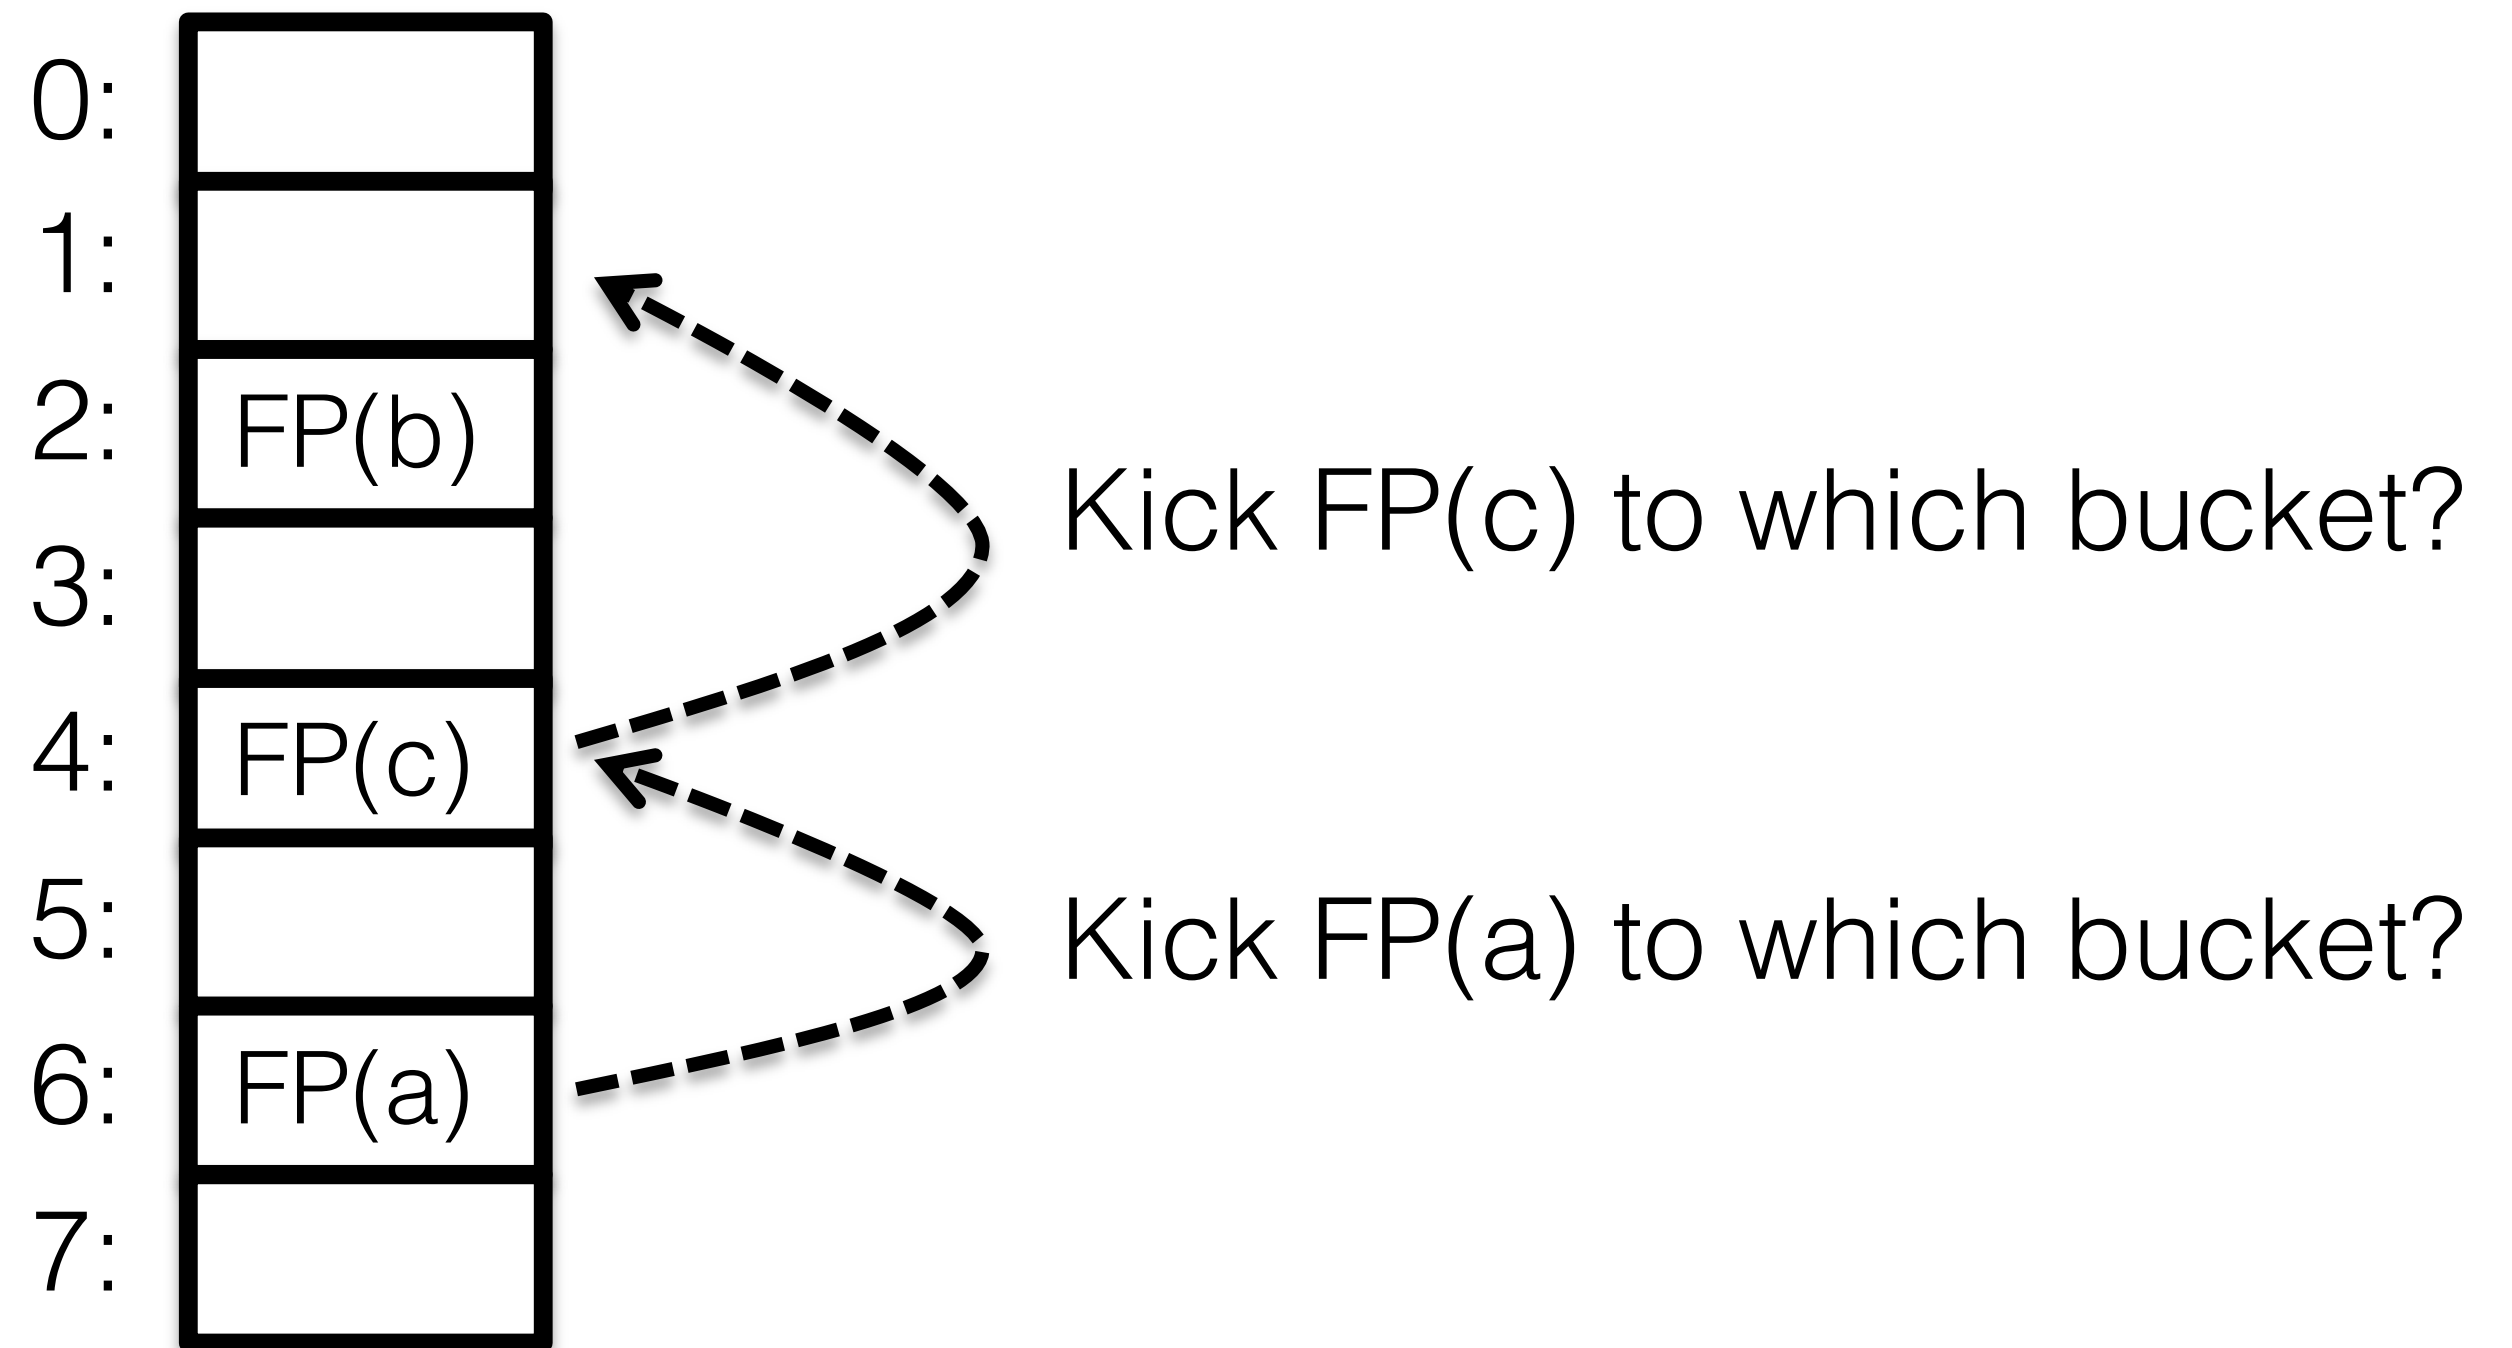
Item



Item



- Cuckoo hashing requires rehashing and displacing items



With only fingerprint,
how to calculate item's alternate bucket?

- Standard Cuckoo Hashing: **two independent hash functions for two buckets**

$$\text{bucket1} = \text{hash}_1(x)$$

$$\text{bucket2} = \text{hash}_2(x)$$

- Partial-key Cuckoo Hashing: **use one bucket and fingerprint to derive the other** [Fan2013]

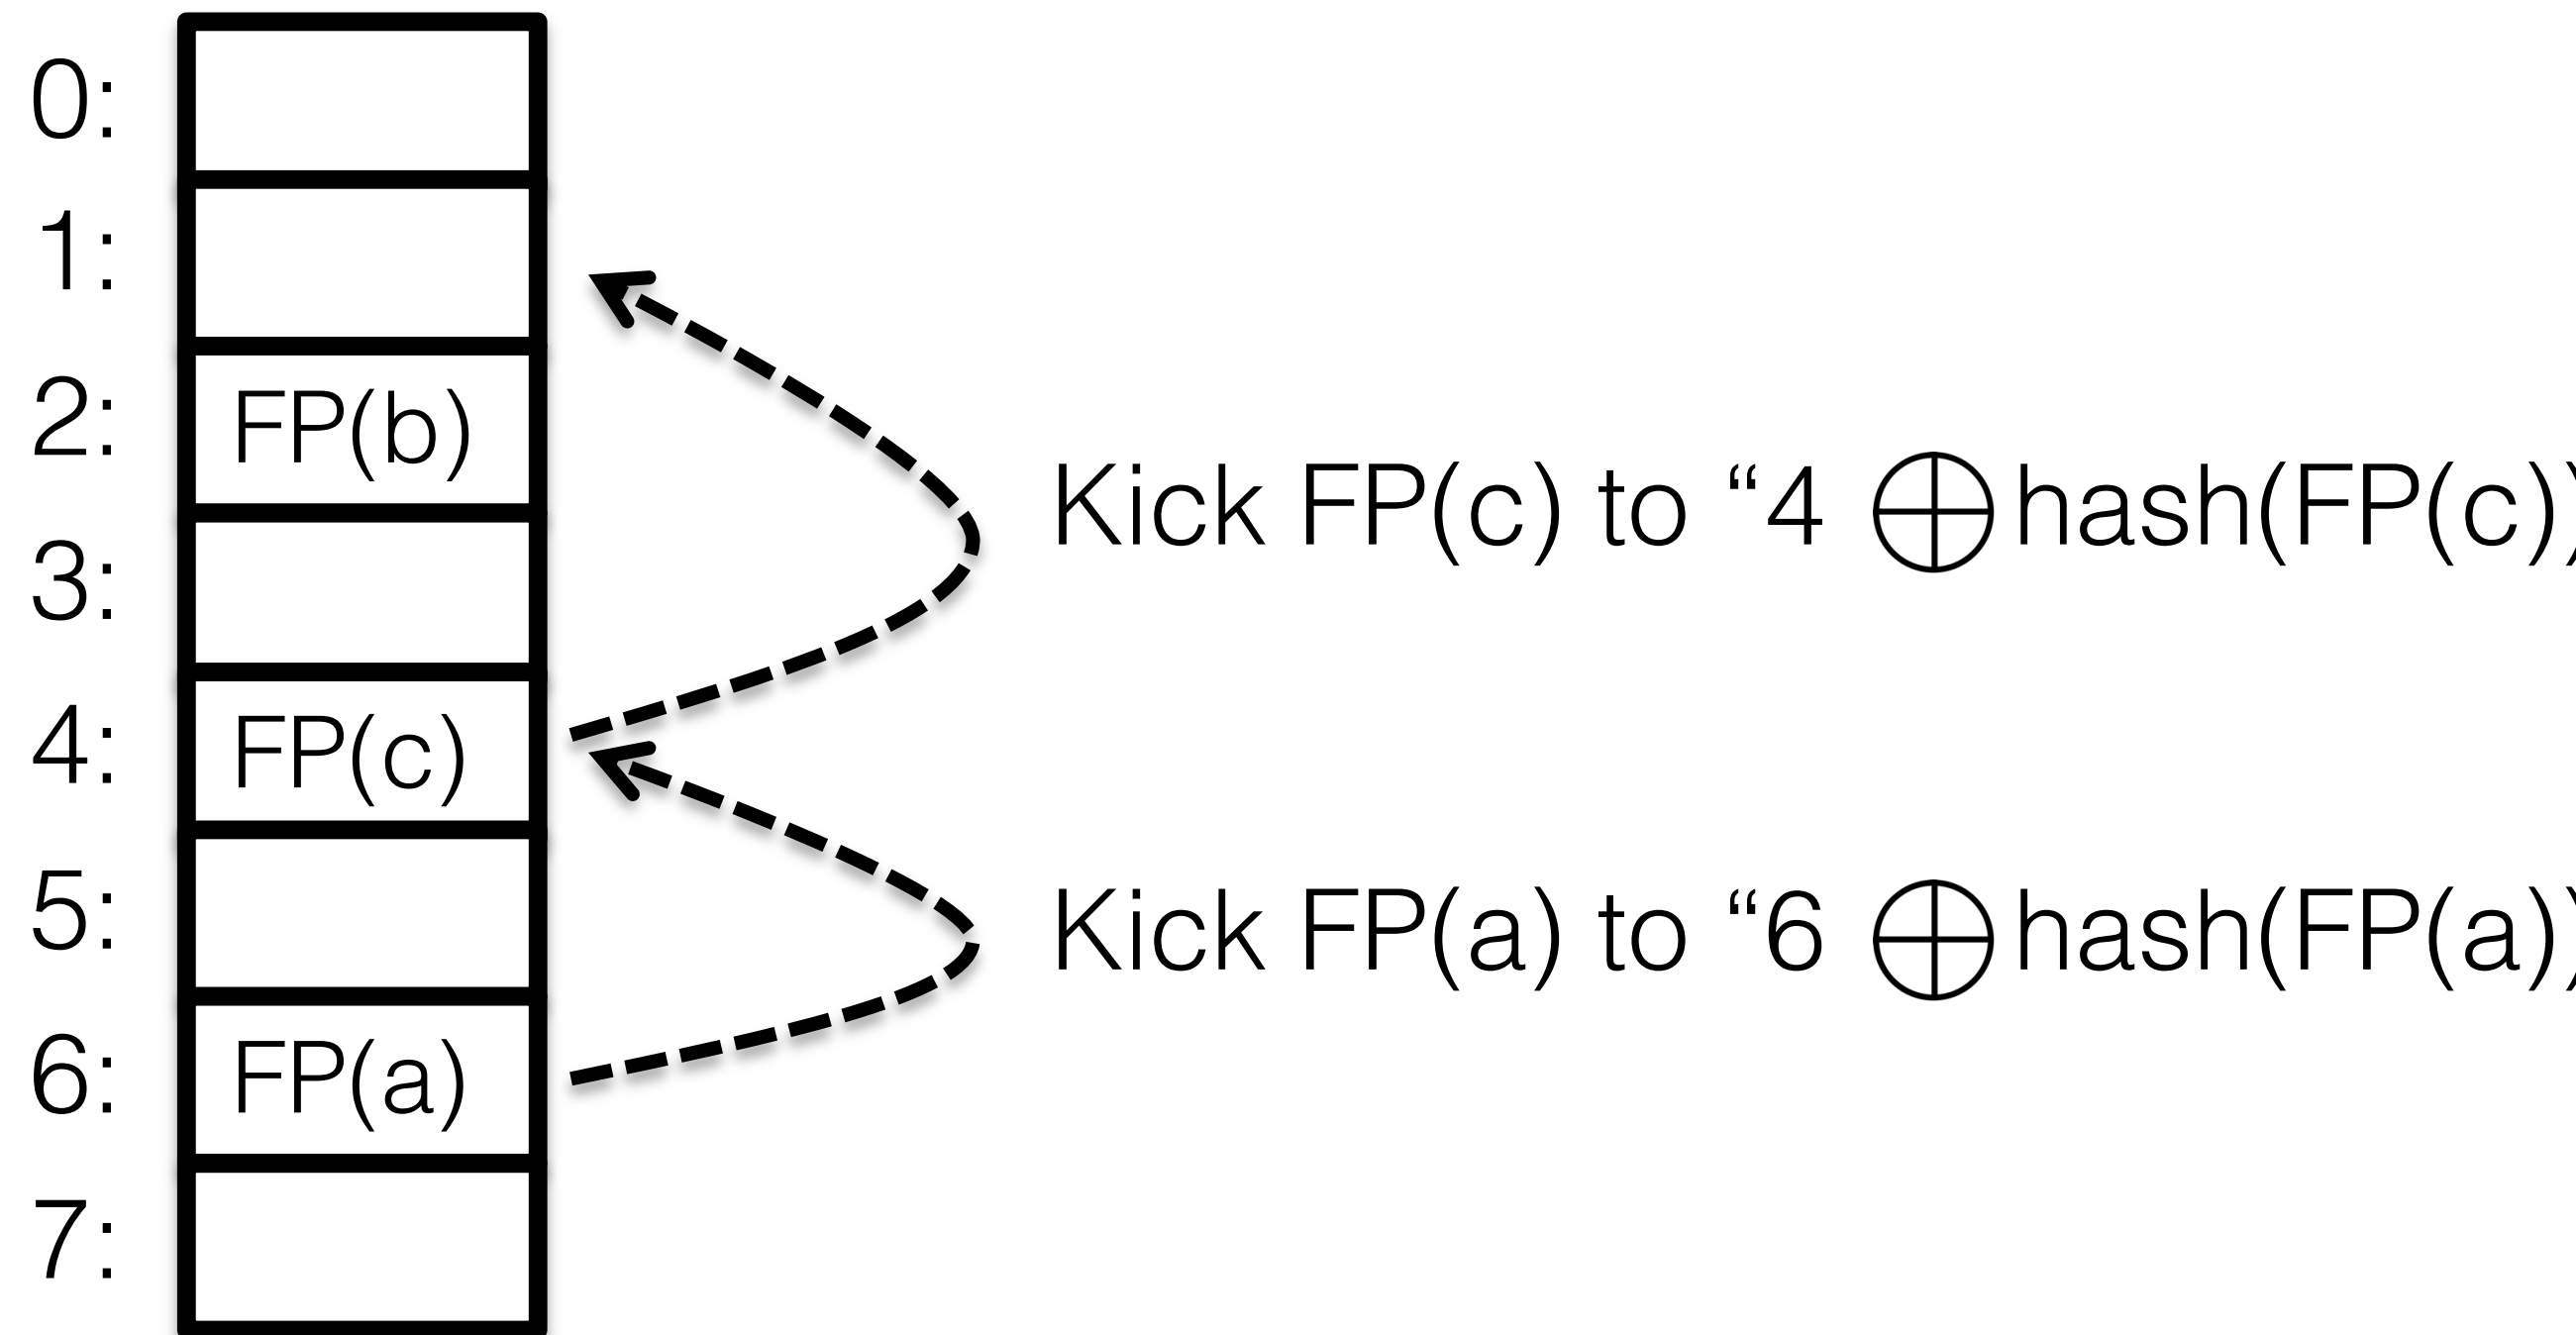
$$\text{bucket1} = \text{hash}(x)$$

$$\text{bucket2} = \text{bucket1} \oplus \text{hash}(\text{FP}(x))$$

To displace existing fingerprint:

$$\text{alternate}(x) = \text{current}(x) \oplus \text{hash}(\text{FP}(x))$$

- Perform cuckoo hashing on fingerprint



Can we still achieve high space utilization with partial-key cuckoo hashing?

Algorithm 1: Insert (x)

$f = \text{fingerprint}(x);$

$i_1 = \text{hash}(x);$

$i_2 = i_1 \oplus \text{hash}(f);$

if bucket[i_1] or bucket[i_2] has an empty entry **then**

 add f to that bucket;

return Done;

// must relocate existing items;

$i = \text{randomly pick } i_1 \text{ or } i_2;$

for $n = 0; n < \text{MaxNumKicks}; n++$ **do**

 randomly select an entry e from bucket[i];

 swap f and the fingerprint stored in entry e ;

$i = i \oplus \text{hash}(f);$

if bucket[i] has an empty entry **then**

 add f to bucket[i];

return Done;

// Hashtable is considered full;

return Failure;

Algorithm 2: Lookup (x)

$f = \text{fingerprint}(x);$

$i_1 = \text{hash}(x);$

$i_2 = i_1 \oplus \text{hash}(f);$

if bucket[i_1] or bucket[i_2] has f **then**

return True;

return False;

Algorithm 3: Delete (x)

$f = \text{fingerprint}(x);$

$i_1 = \text{hash}(x);$

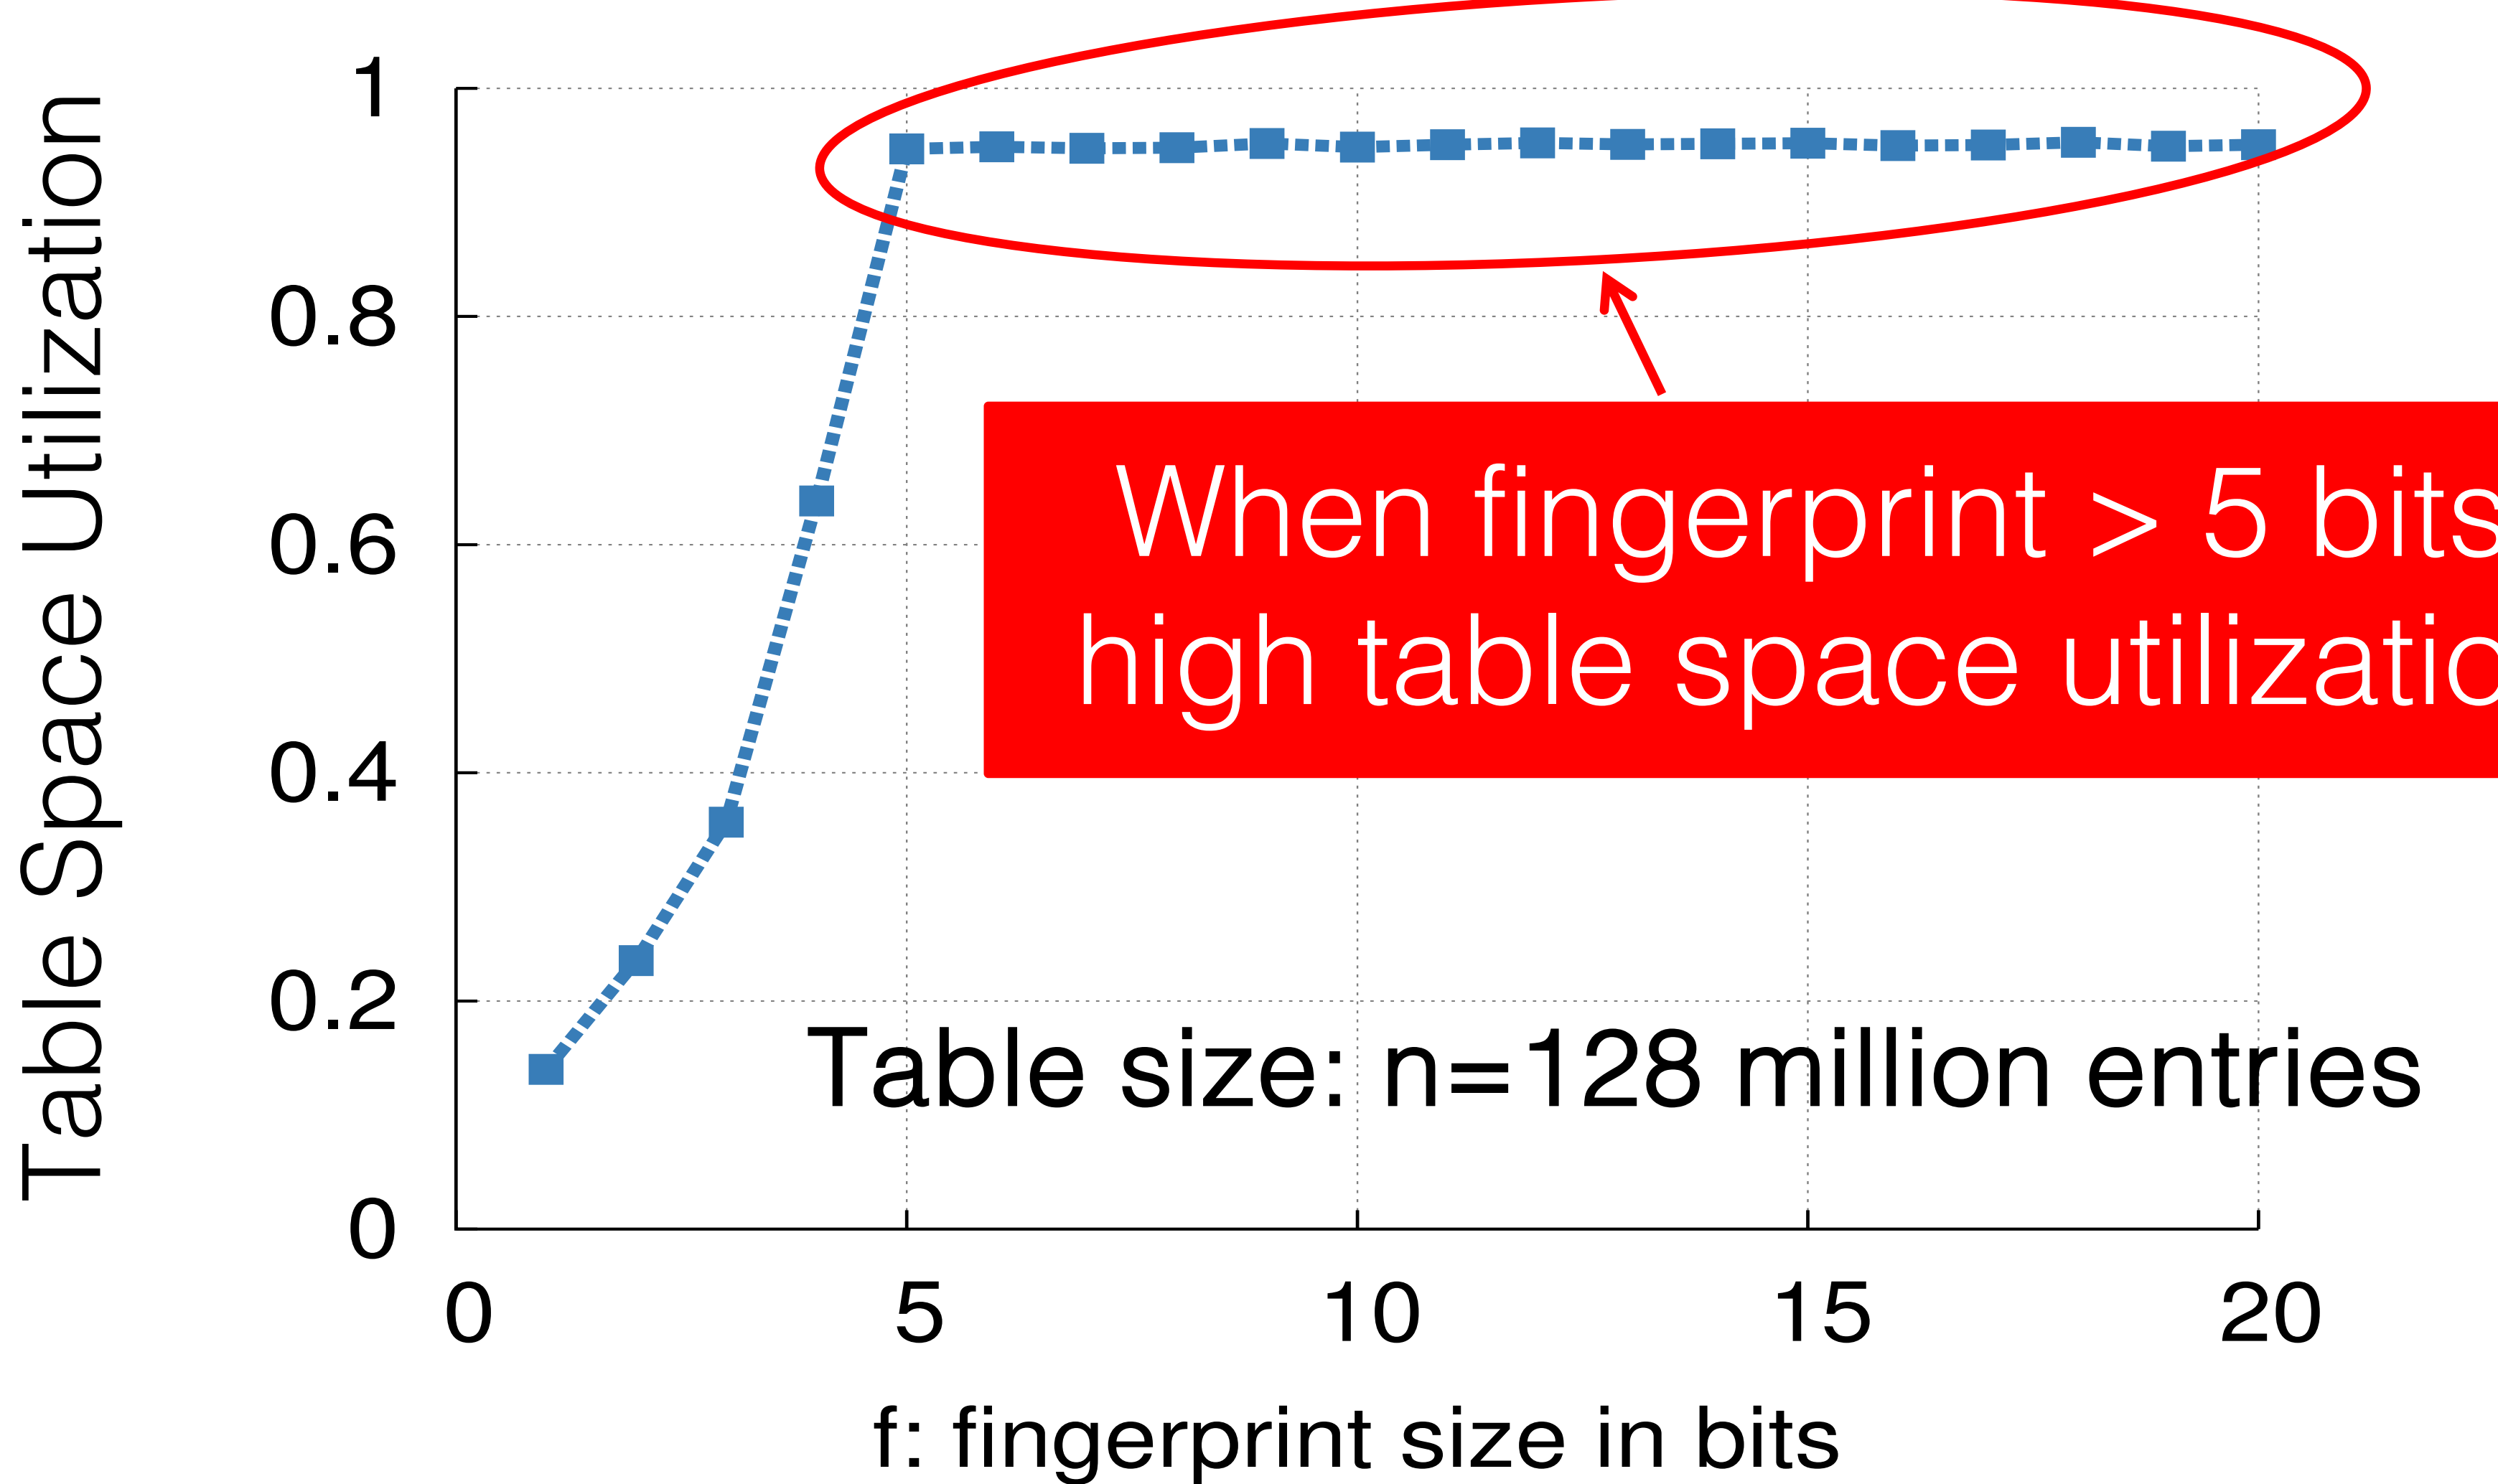
$i_2 = i_1 \oplus \text{hash}(f);$

if bucket[i_1] or bucket[i_2] has f **then**

 remove a copy of f from this bucket;

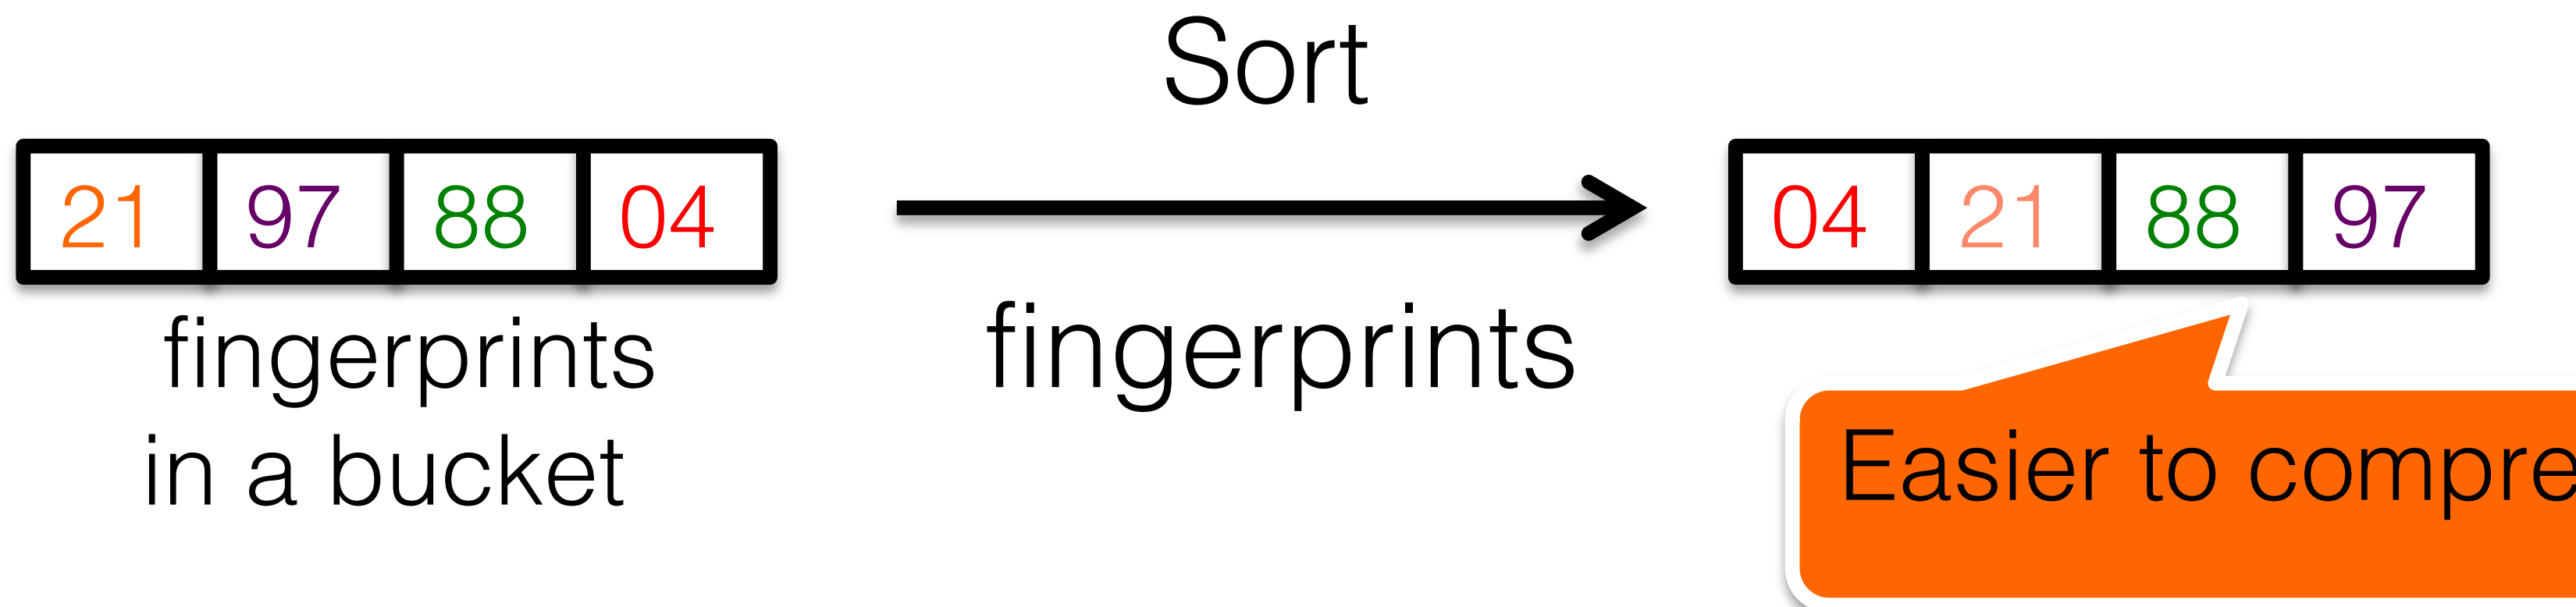
return True;

return False;



- Fingerprint must be $\Omega(\log(n)/b)$ bits in theory
 - n : hash table size, b : bucket size

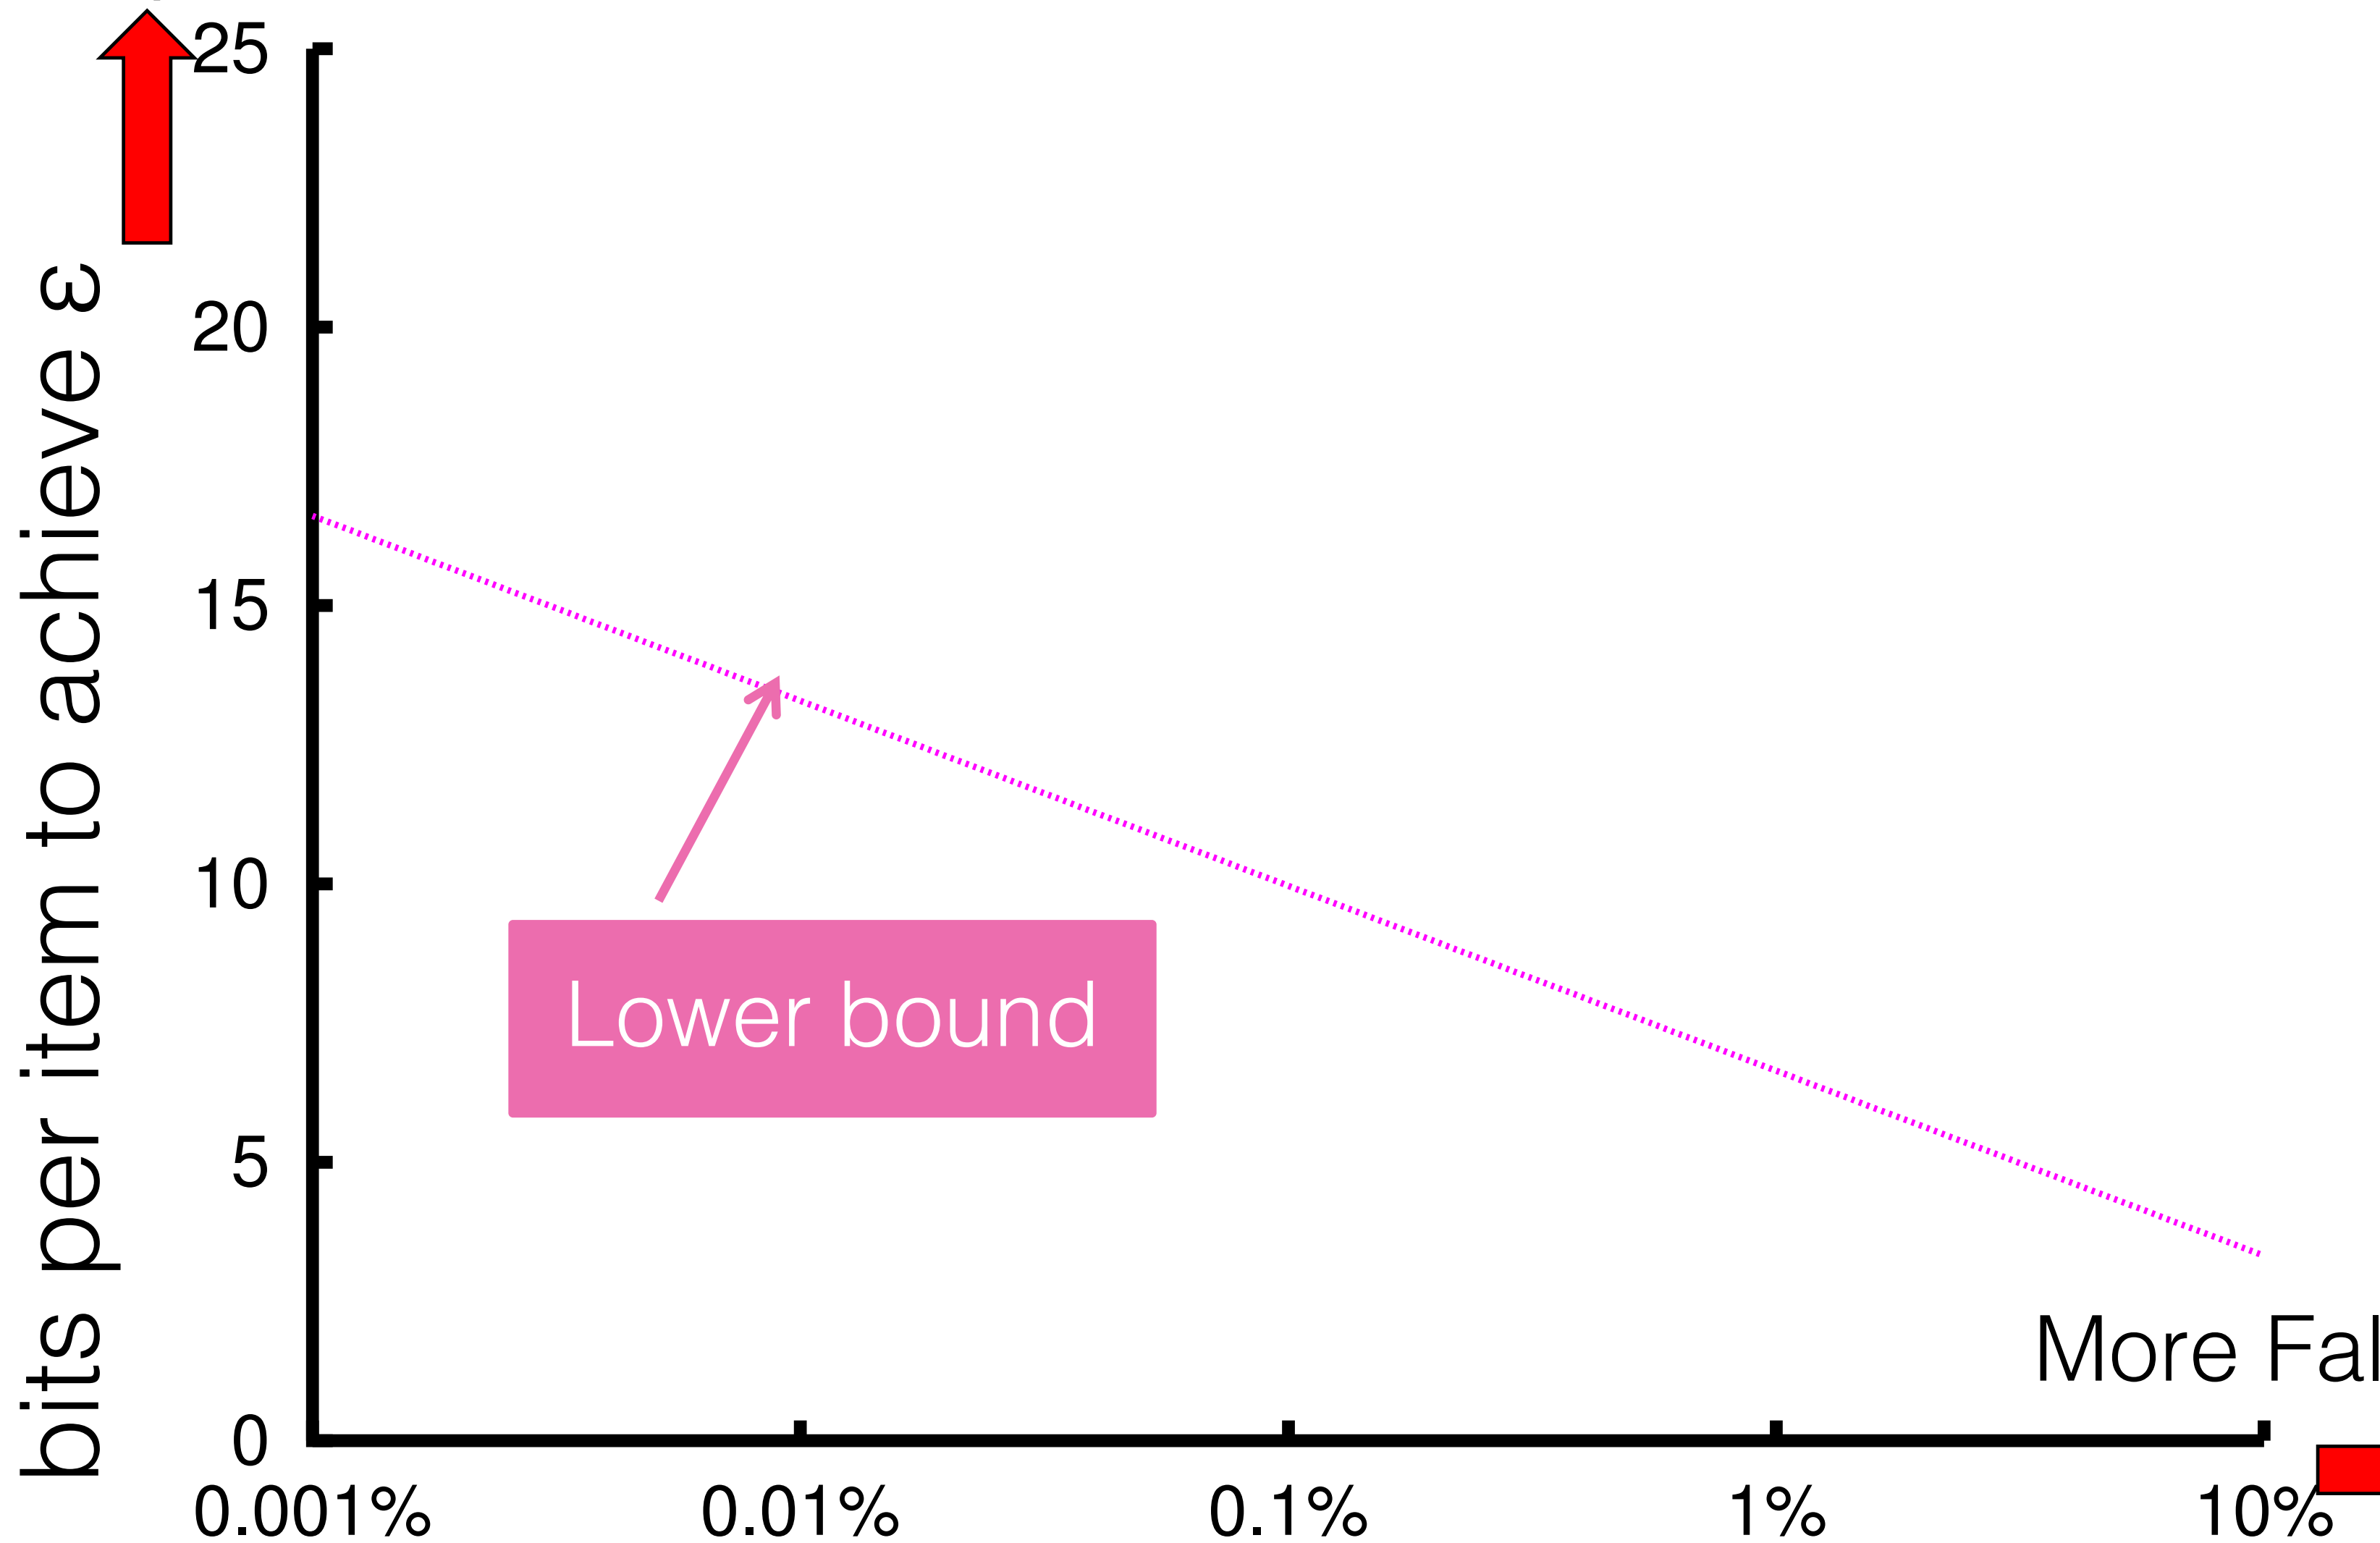
- Based on observation:
 - A monotonic sequence of integers is easier to compress [\[Bonomi2006\]](#)
- Semi-Sorting:
 - Sort fingerprints sorted in each bucket
 - Compress sorted fingerprints



- + For 4-way bucket, save one bit per item
- Slower lookup / insert

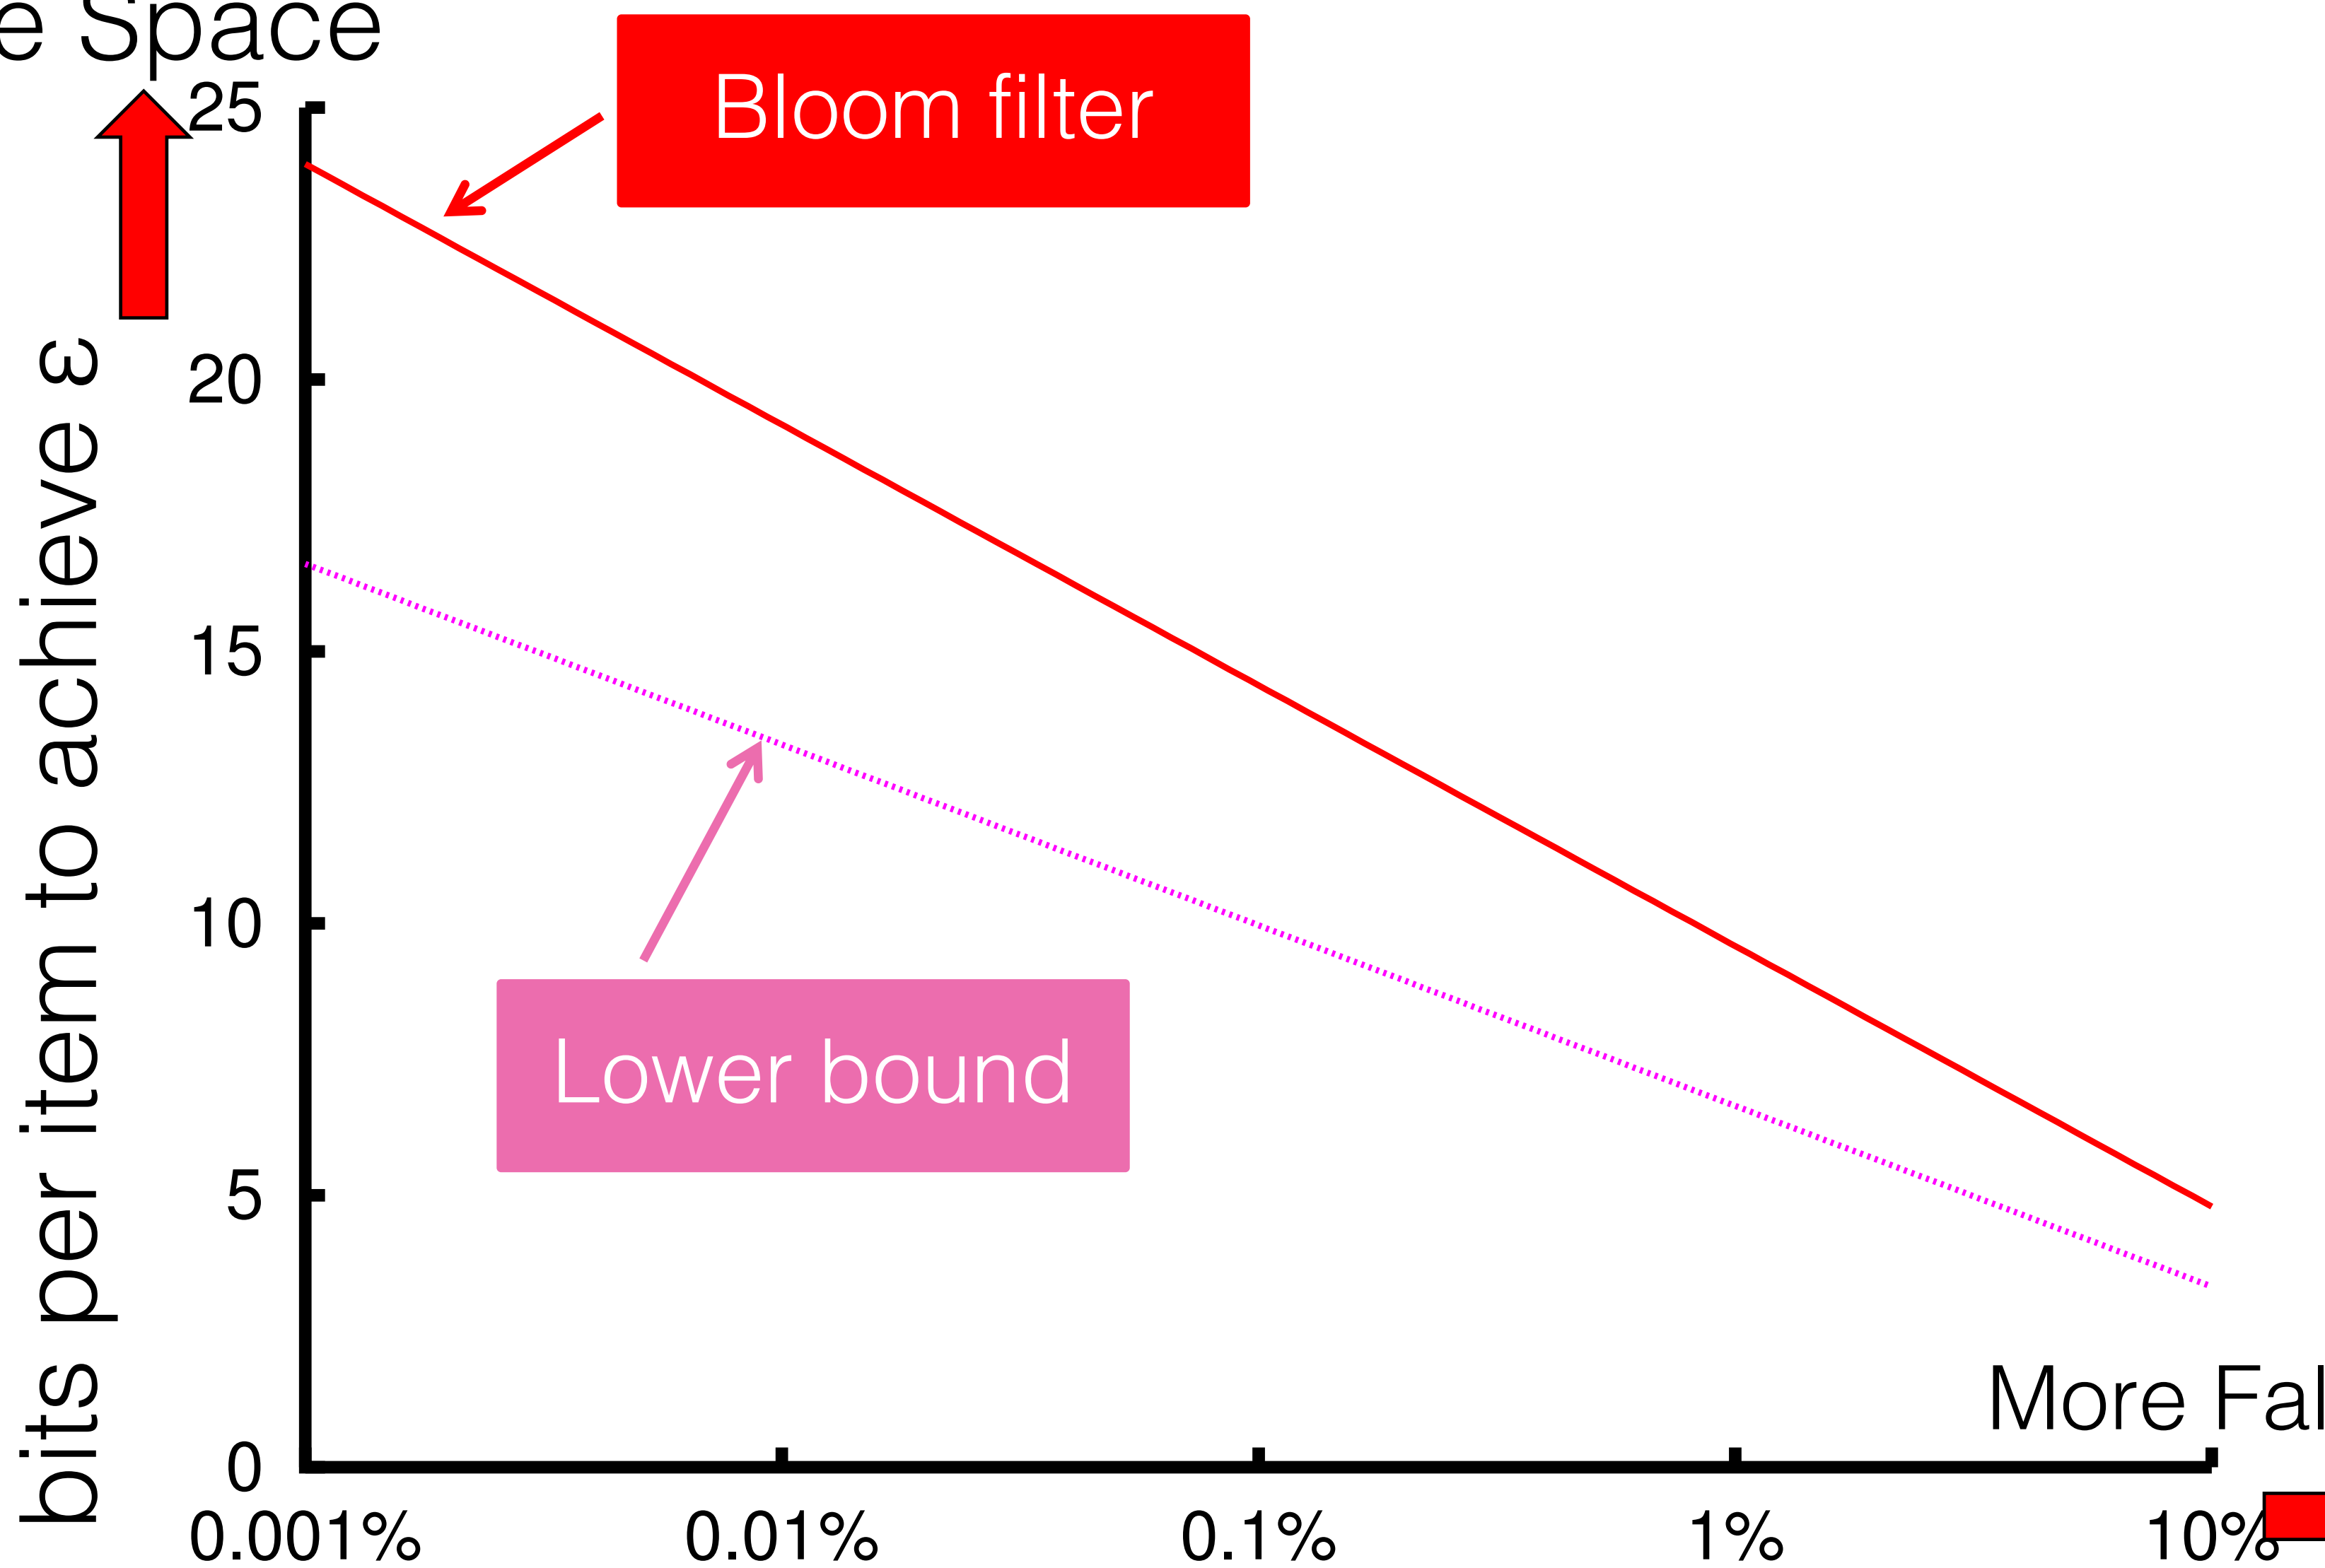
[\[Bonomi2006\]](#) Beyond Bloom filters: From approximate membership checks to approximate state machines.

More Space



ϵ : target false positive rate

More Space



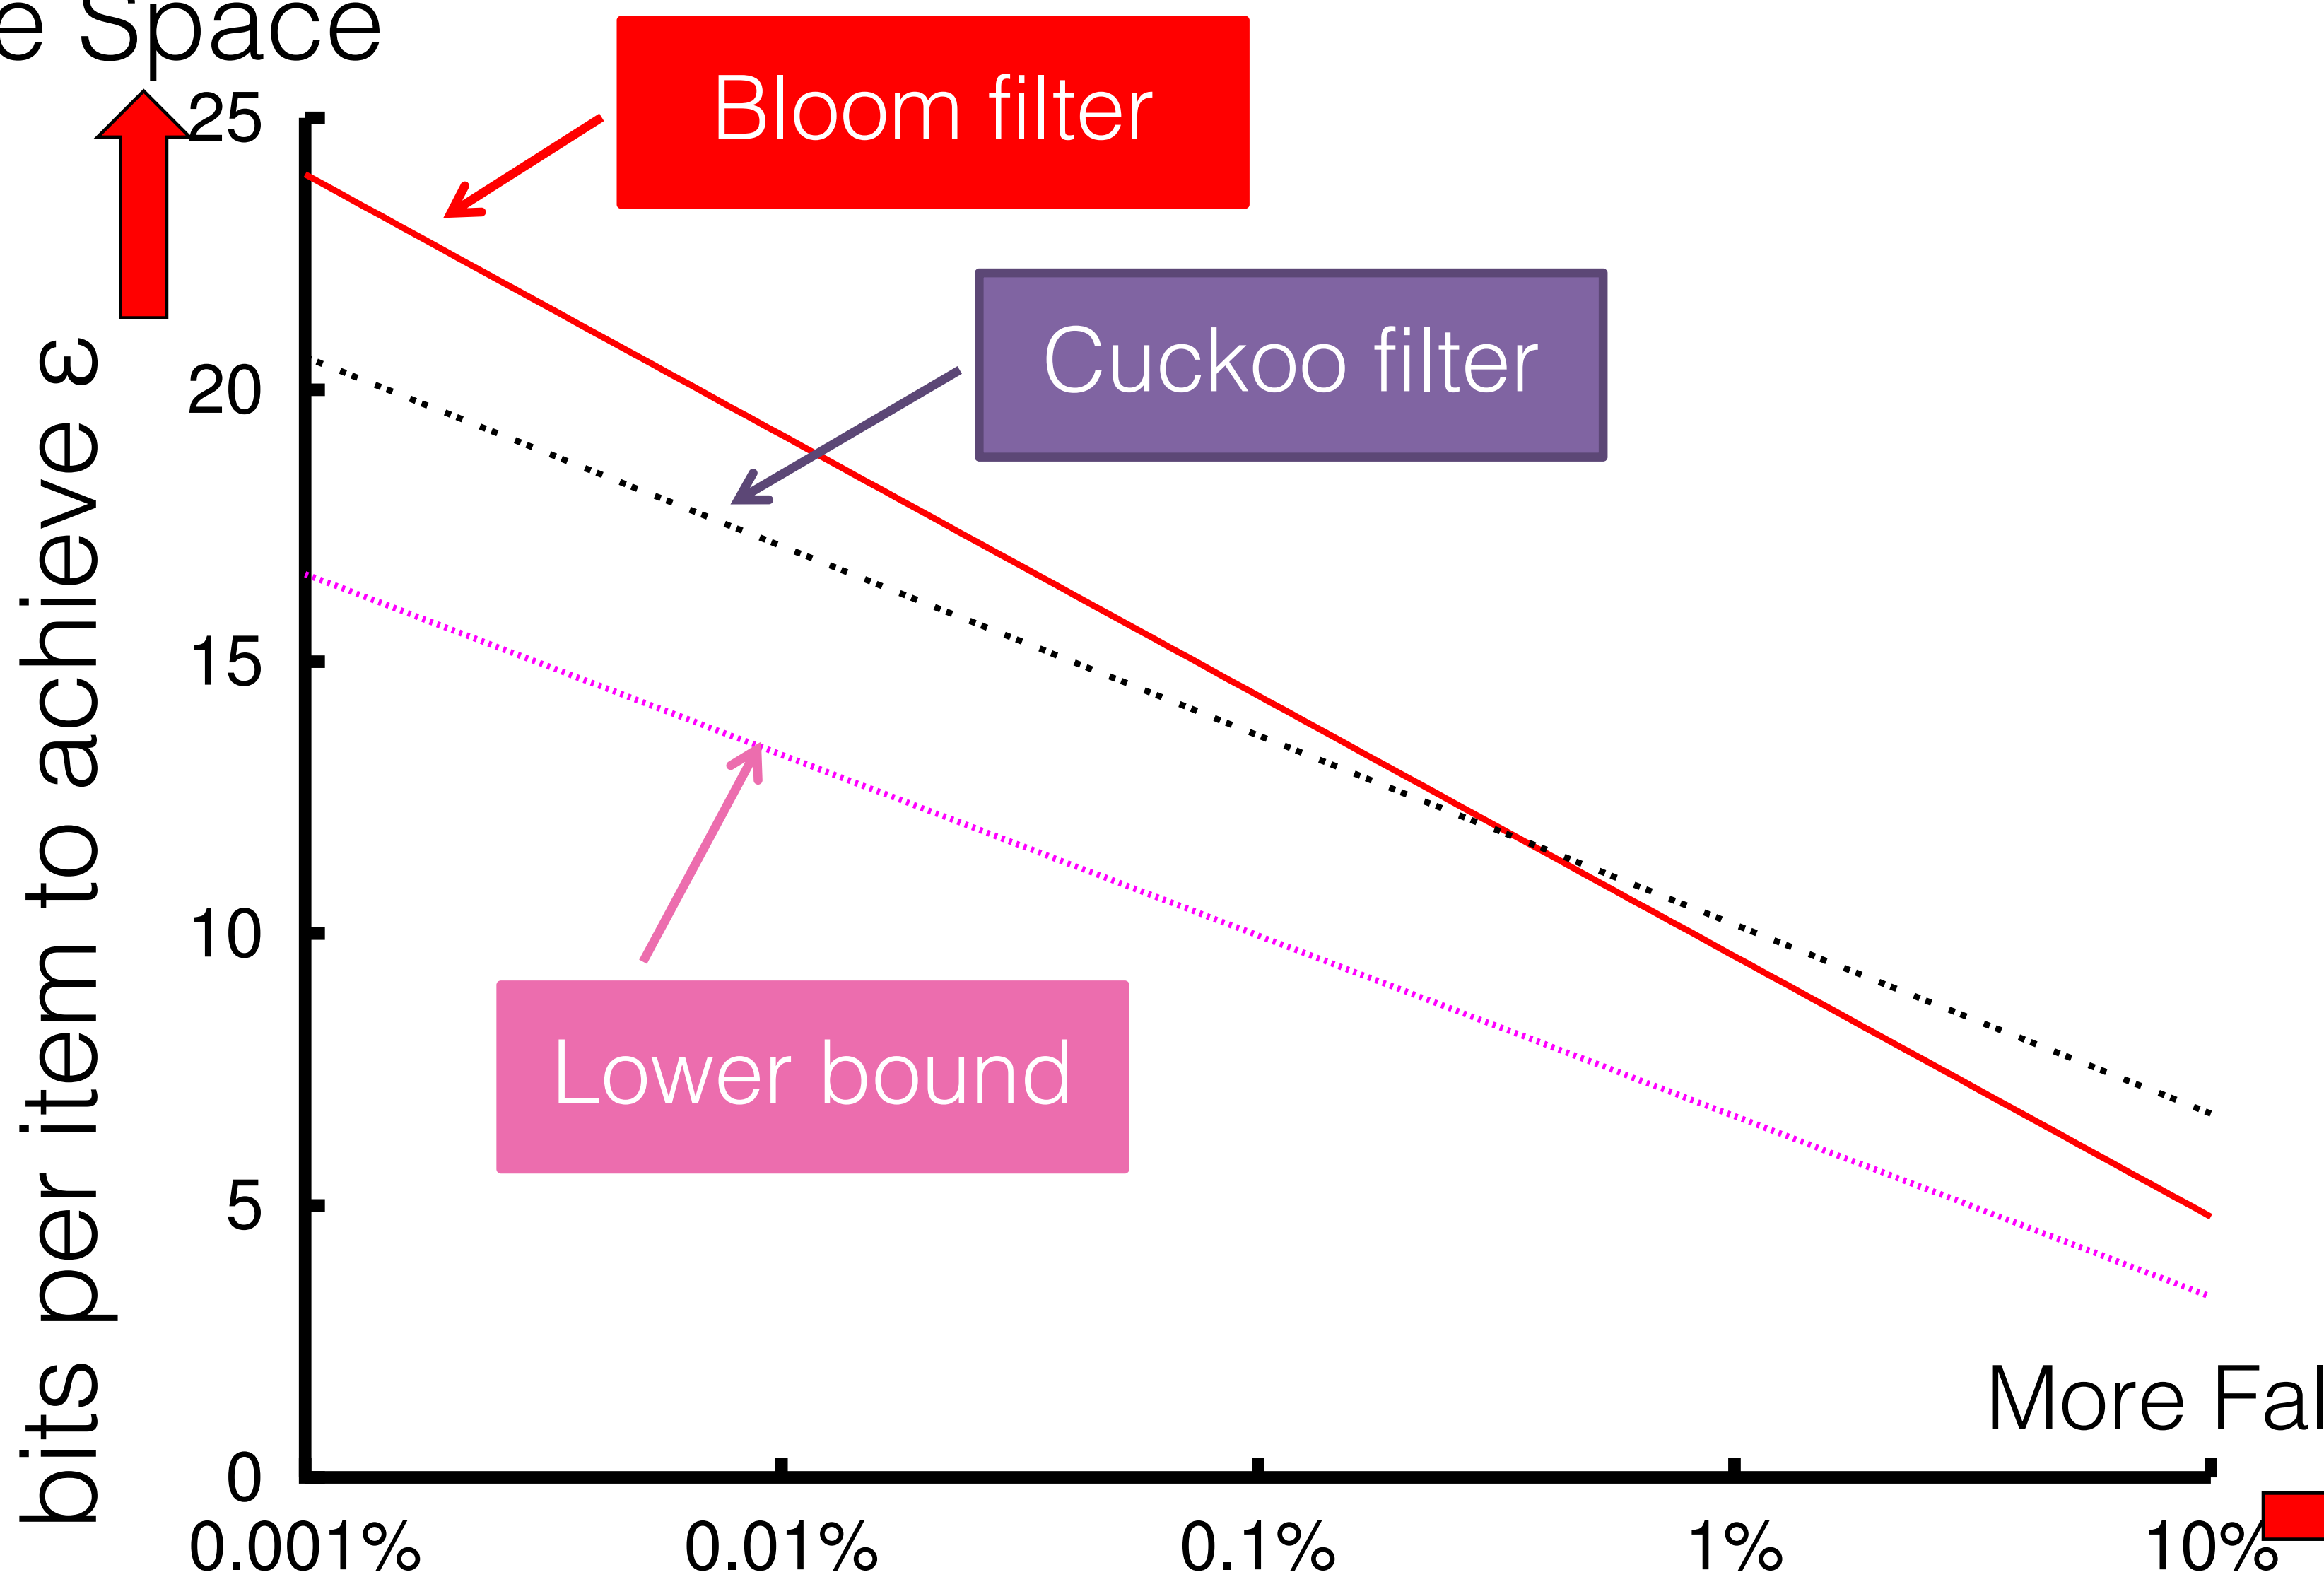
Lower bound

Bloom filter

More False

ϵ : target false positive rate

More Space



ϵ : target false positive rate

More Space

bits per item to achieve ϵ

25
20
15
10
5
0

0.001%

0.01%

0.1%

1%

10%

Bloom filter

Cuckoo filter

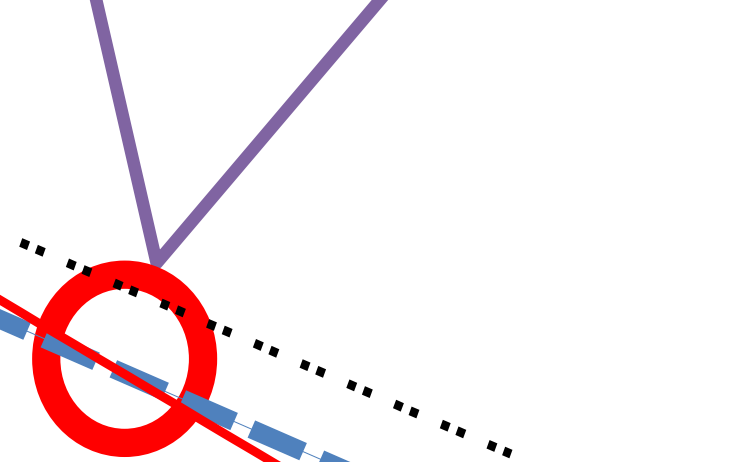
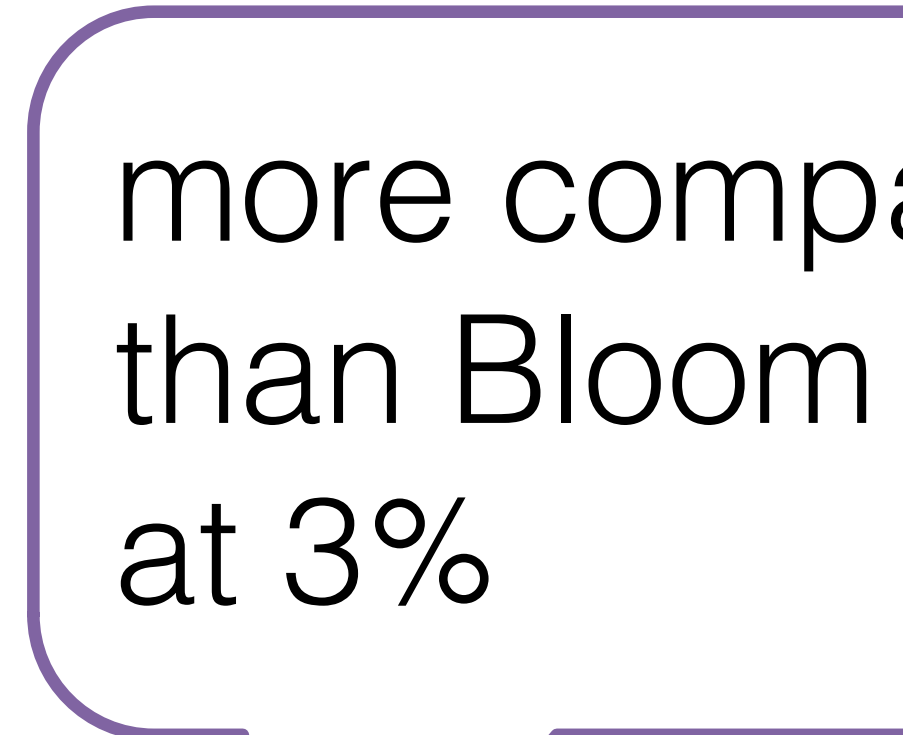
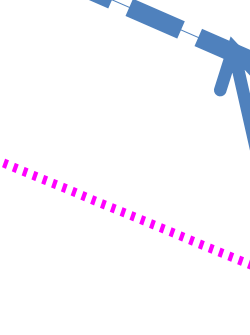
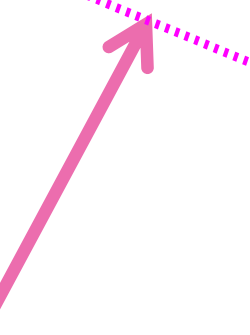
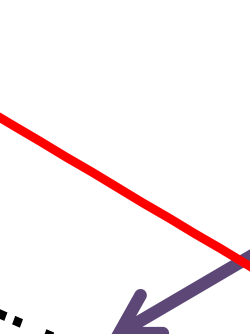
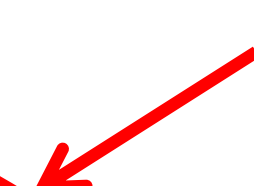
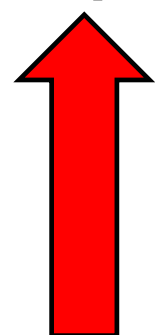
Lower bound

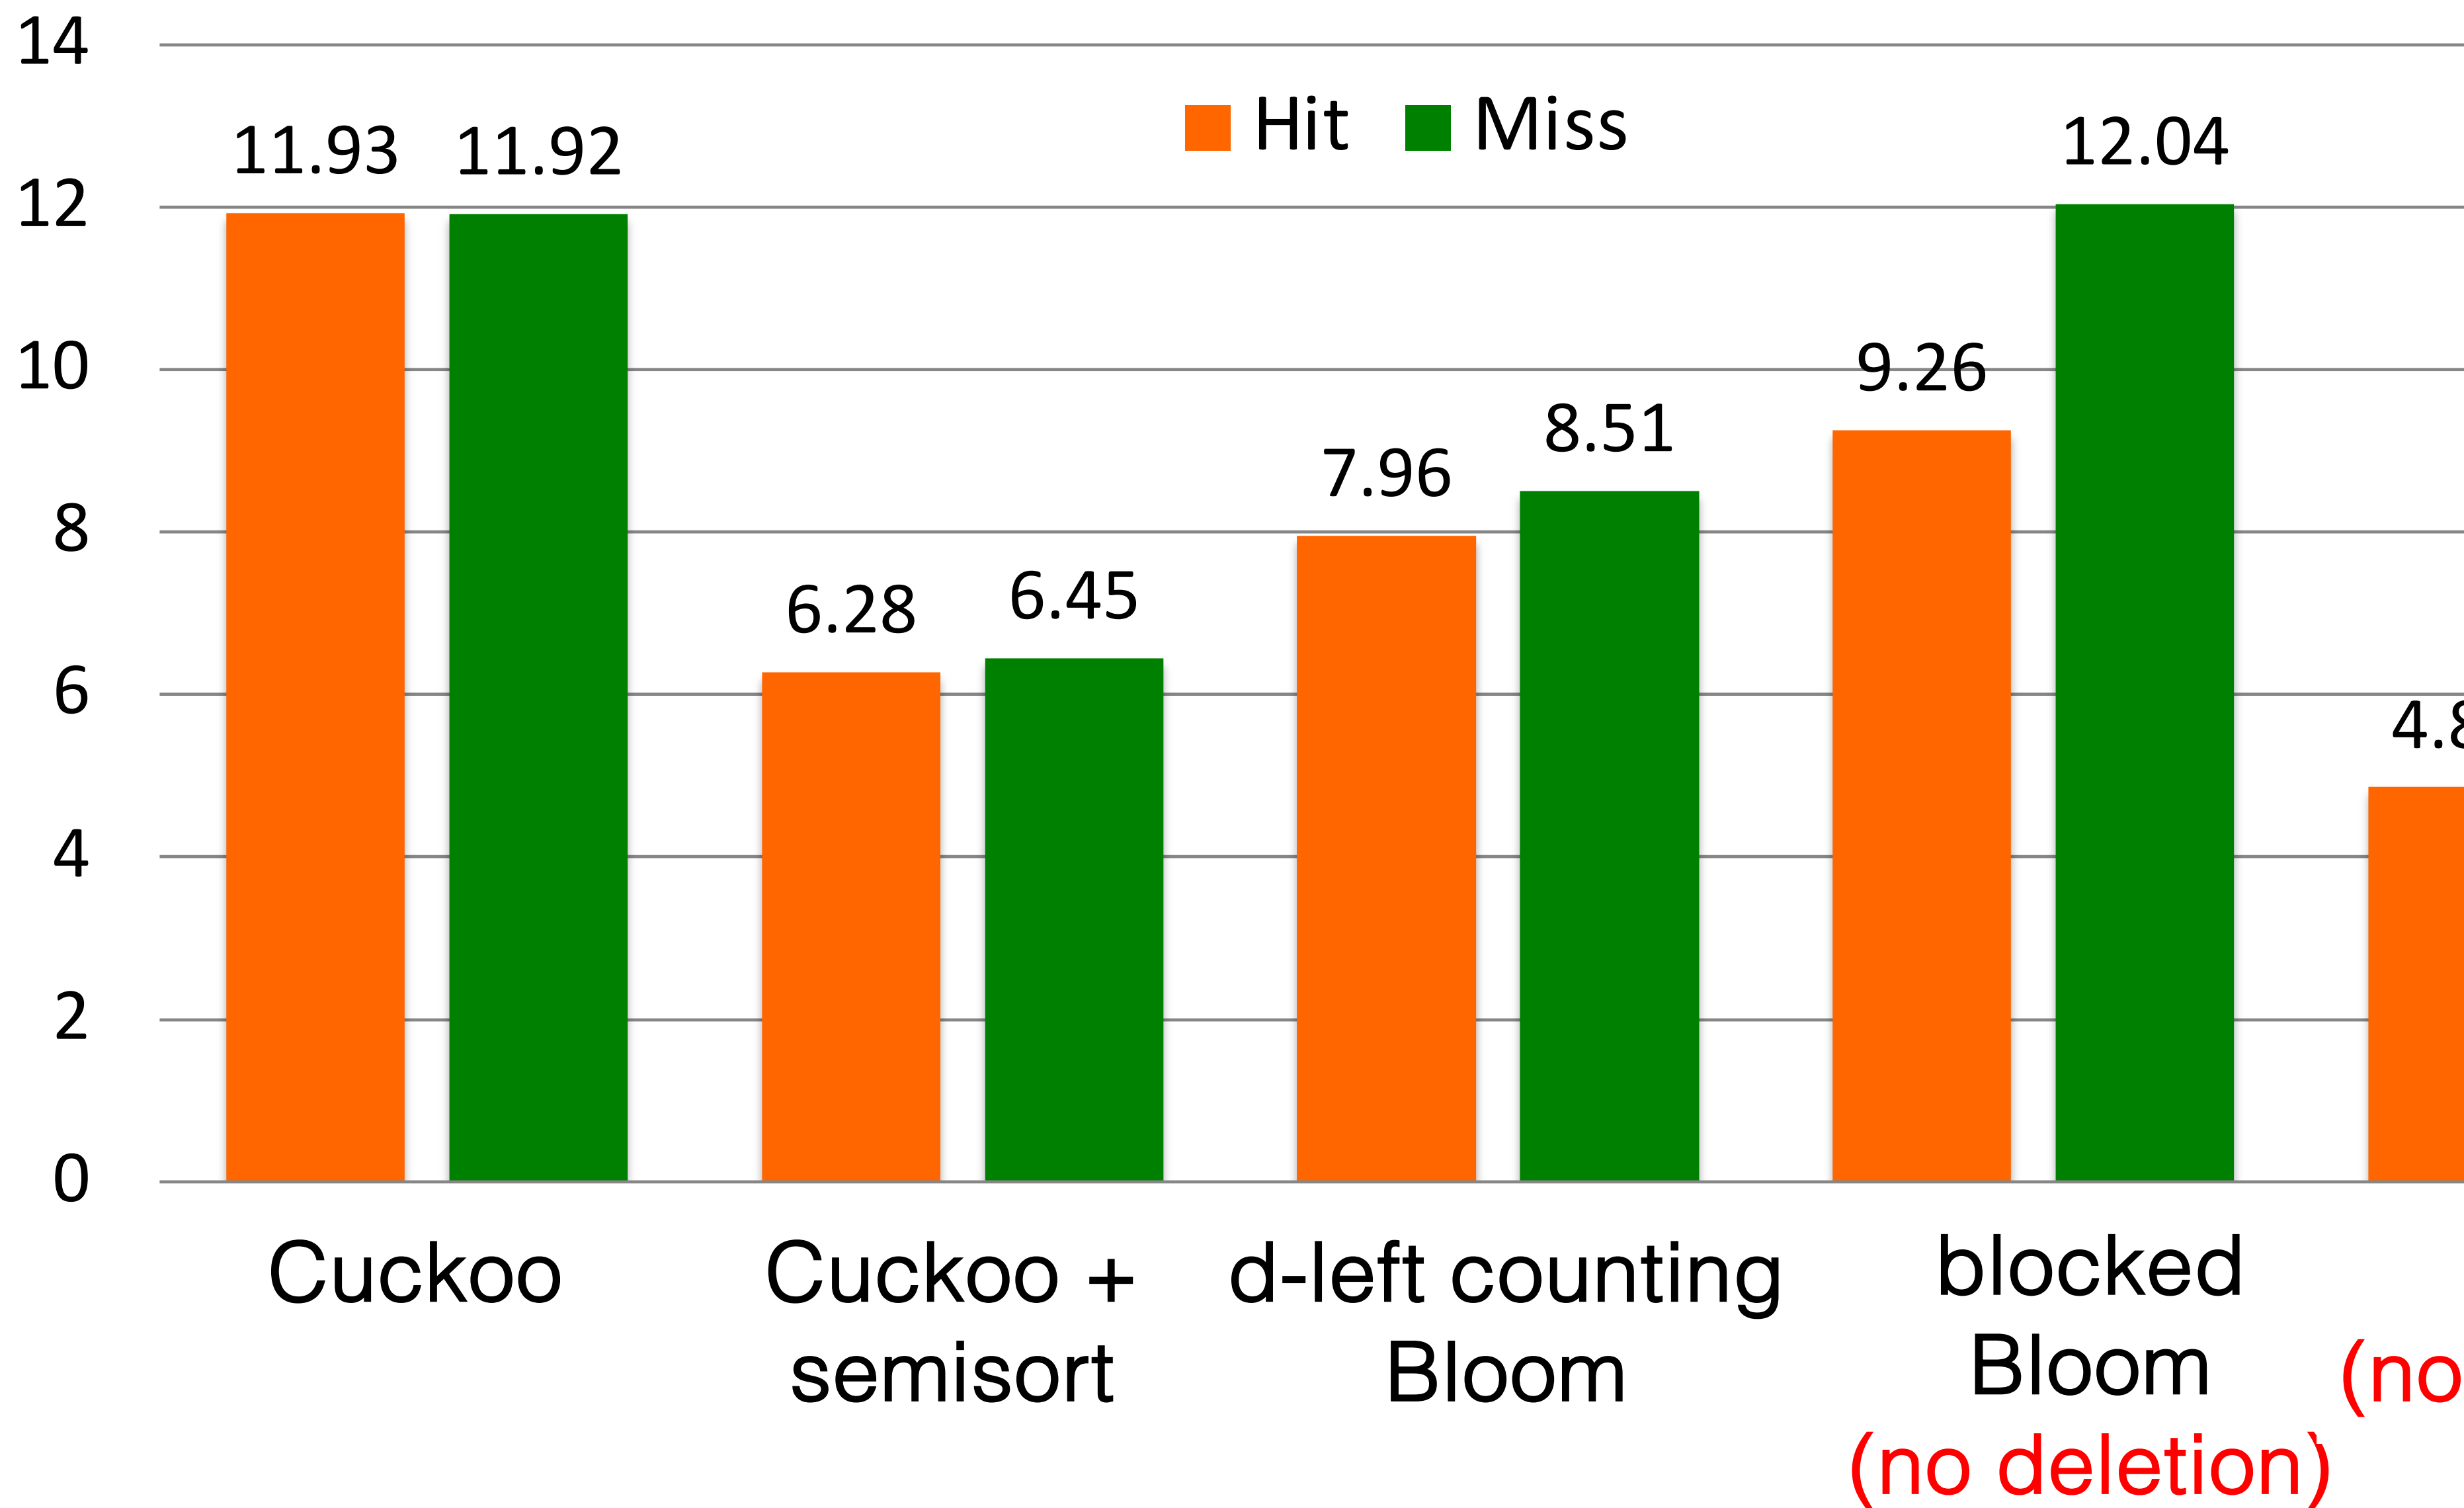
Cuckoo filter +
semisorting

more compact
than Bloom
at 3%

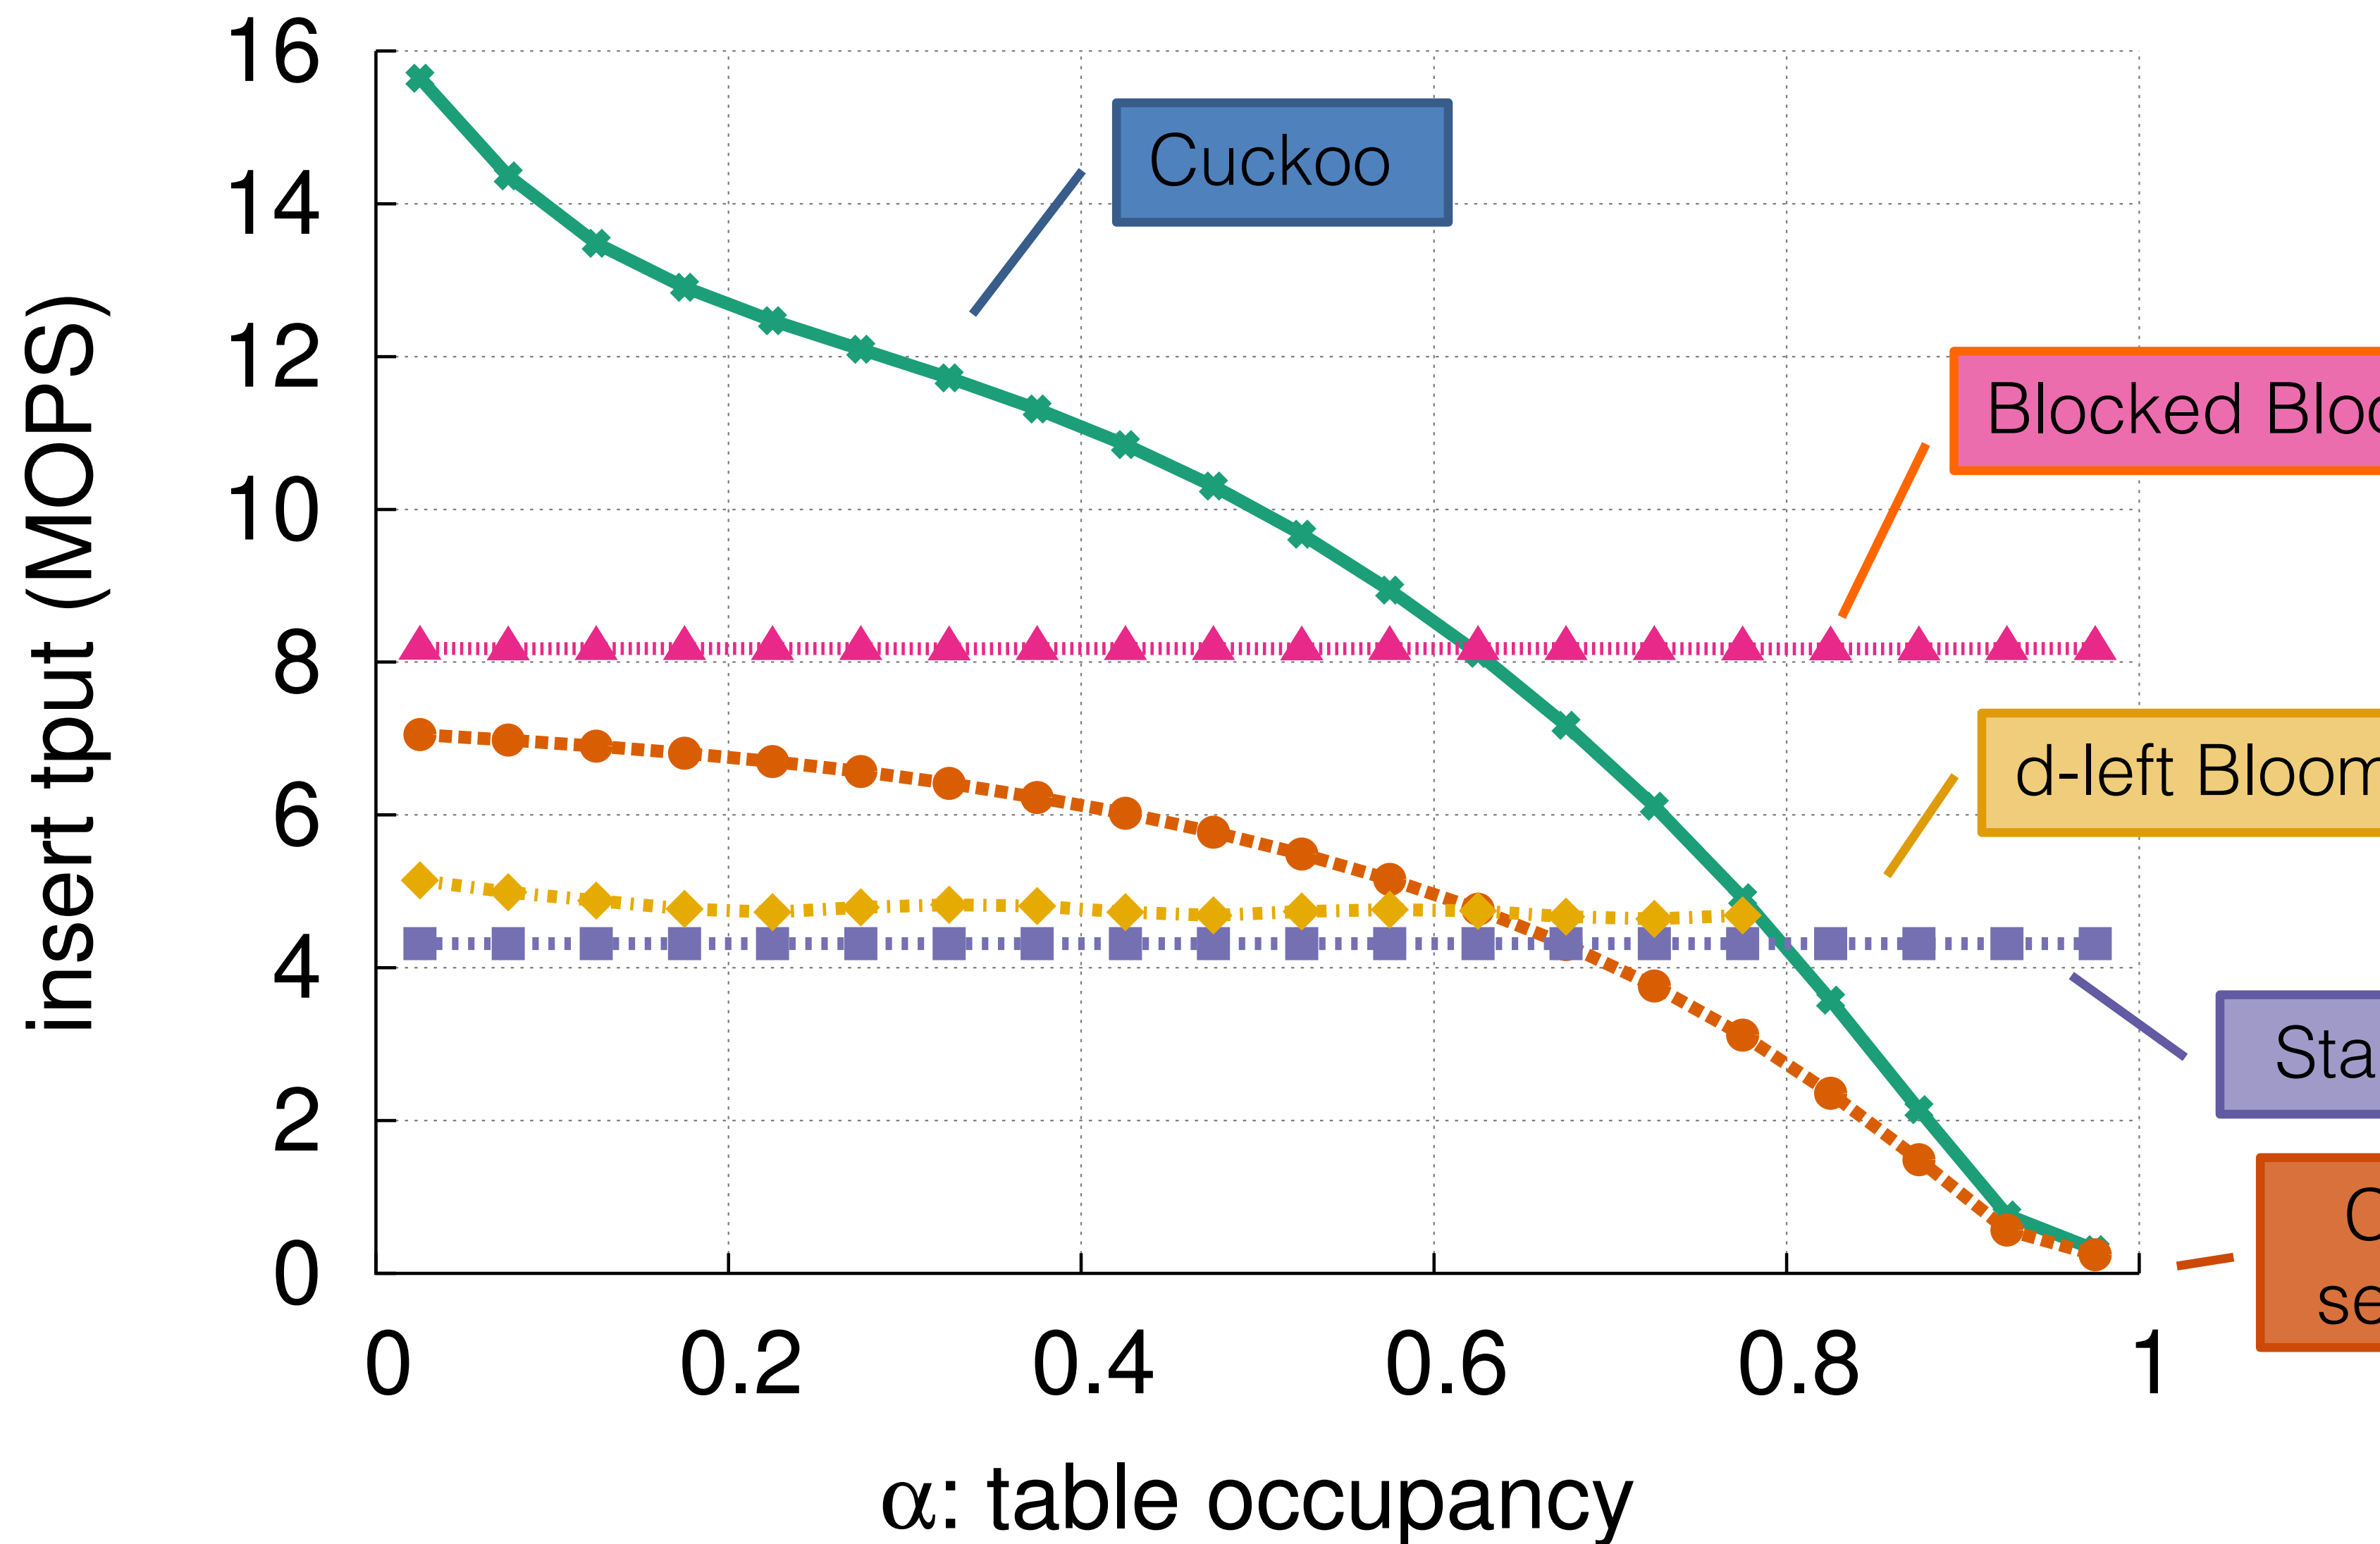
More False

ϵ : target false positive rate





Cuckoo filter is among the fastest regardless workload



Cuckoo filter has decreasing insert rate, but overall is only slower than blocked Bloom filter.

References:

- Mining massive Datasets by Leskovec, Rajaraman, Ullman, Chapter 4.
- Primary reference for this lecture
 - Survey on Bloom Filter, Broder and Mitzenmacher 2005,
<https://www.eecs.harvard.edu/~michaelm/postscripts/im2005b.pdf>
- Others
 - Randomized Algorithms by Mitzenmacher and Upfal.
- Cuckoo filter: Fan, Bin, Dave G. Andersen, Michael Kaminsky, and Michael D. Mitzenmacher. "Cuckoo filter: Practically better than bloom." In Proceedings the 10th ACM International on Conference on emerging Networking Experiments and Technologies, pp. 75-88. 2014.